

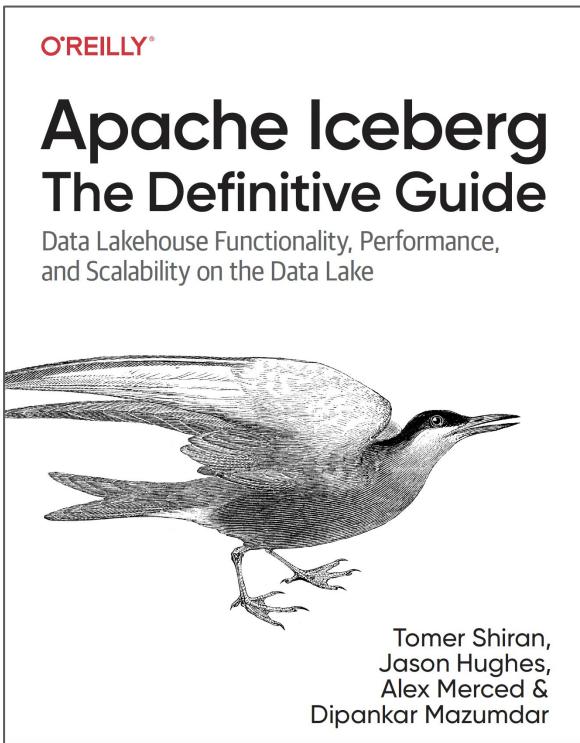
# Apache Iceberg



An Architectural Look Under the Covers

Community Over Code, North America, 2023

# About Me



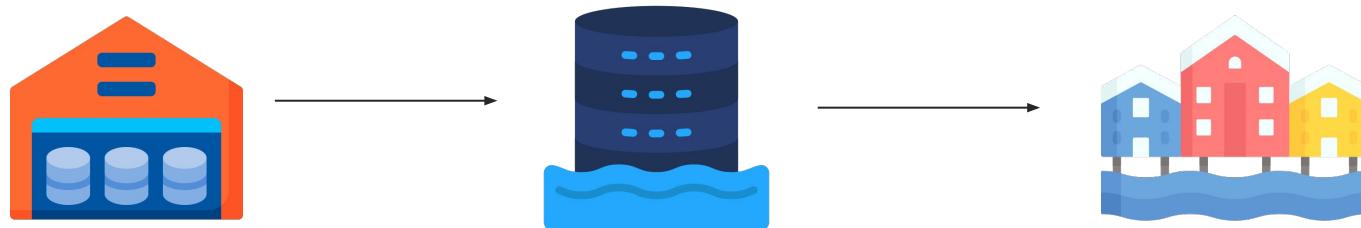
- **Current:** Data (Eng/Sci) Advocate at Dremio
- **Open source:** Apache Iceberg, Arrow, Project Nessie
- **Past:** BI, Data Viz, ML

# Key Takeaways

- Evolution of data architecture
- Data Lakehouse
- Architectural deep dive: Iceberg
- How queries work under covers?
- Design Benefits

# Evolution of Data Architecture

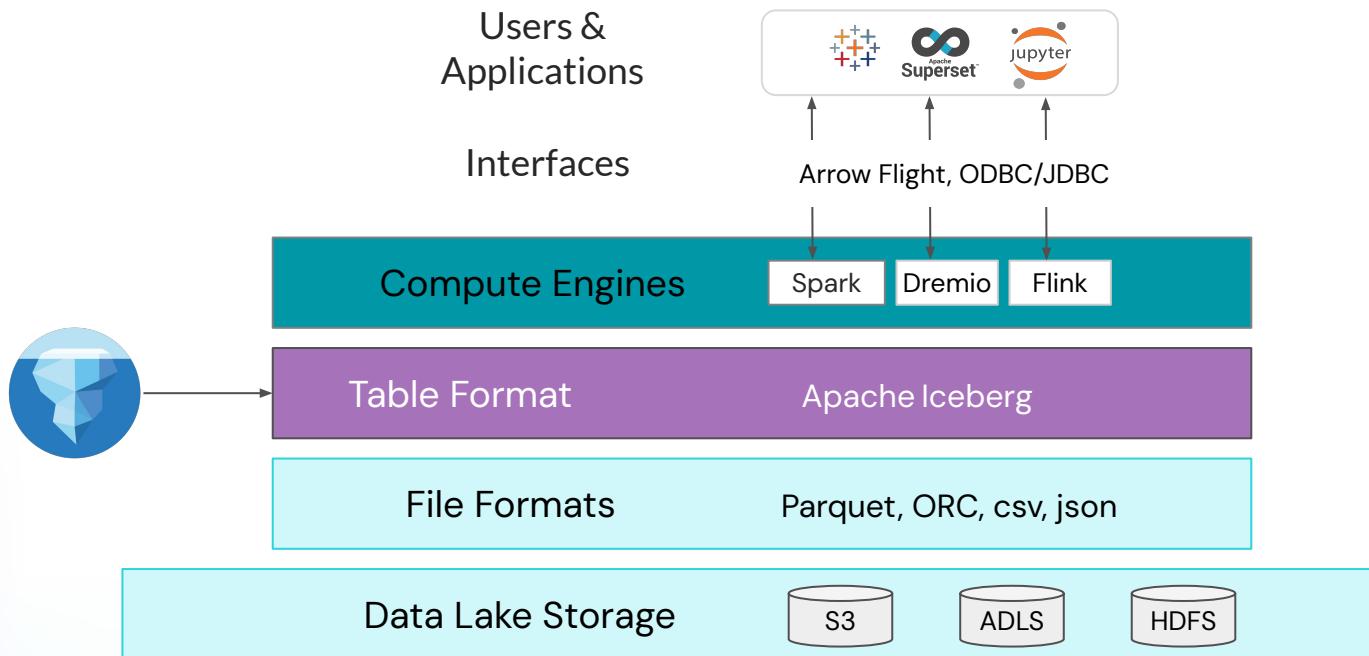
How did we get here?



Centralized, reliable data platform  
Democratize data

Data Warehouse → Data Lakes → Lakehouse

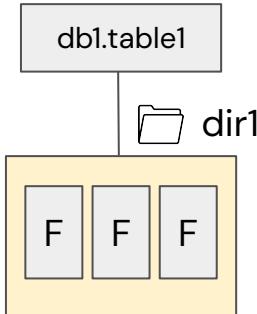
# Iceberg in a Data Lakehouse



# What is a table format?

- A way to organize a dataset's files to present them as a single "table"
- A way to answer the question "what data is in this table?"

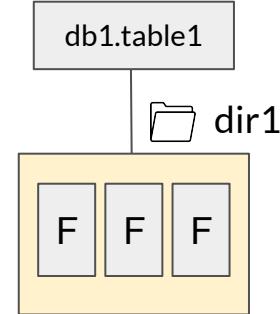
Old way  
(Hive)



A table's contents is all files  
in that table's directories

# Hive table format

- A table's contents is all files in that table's directories
- The old de-facto standard



## Pros

- Works with basically every engine since it's been the de-facto standard for so long
- More efficient access patterns than full-table scans for every query
- File format agnostic
- Atomically update a whole partition
- Single, central answer to "what data is in this table" for the whole ecosystem

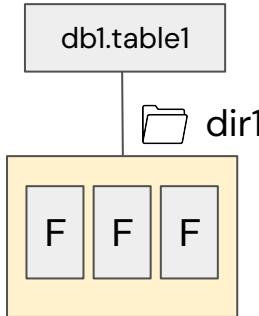
## Cons

- Smaller updates are very inefficient
- No way to change data in multiple partitions safely
- In practice, multiple jobs modifying the same dataset don't do so safely
- All of the directory listings needed for large tables take a *long* time
- Users have to know the physical layout of the table
- Hive table statistics are often stale

# How can we resolve these issues?

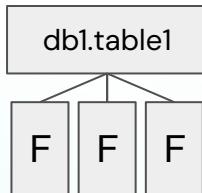
→ We need a new table format

Old way  
(Hive)



A table's contents is all files  
in that table's directories

New way



A table is a **canonical list of files**

## NETFLIX 's goals

- Table correctness/consistency
- Faster query planning and execution
- Allow users to not worry about the physical layout of the data
- Table evolution
- Accomplish all of these at scale

# What Iceberg is and isn't



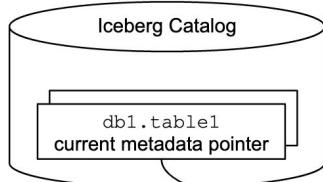
- Table format specification
- A set of APIs and libraries for interaction with that specification
  - These libraries are leveraged in other engines and tools that allow them to interact with Iceberg tables
- A storage engine
- An execution engine
- A service

# Architectural Deep Dive

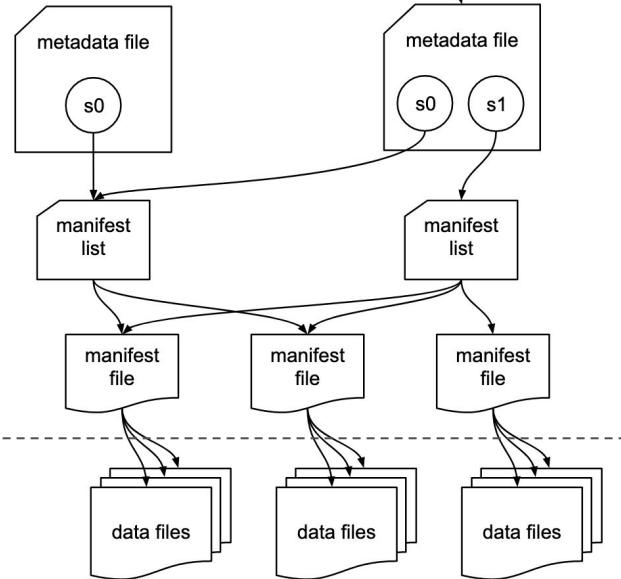
# Iceberg table format



catalog layer



metadata layer

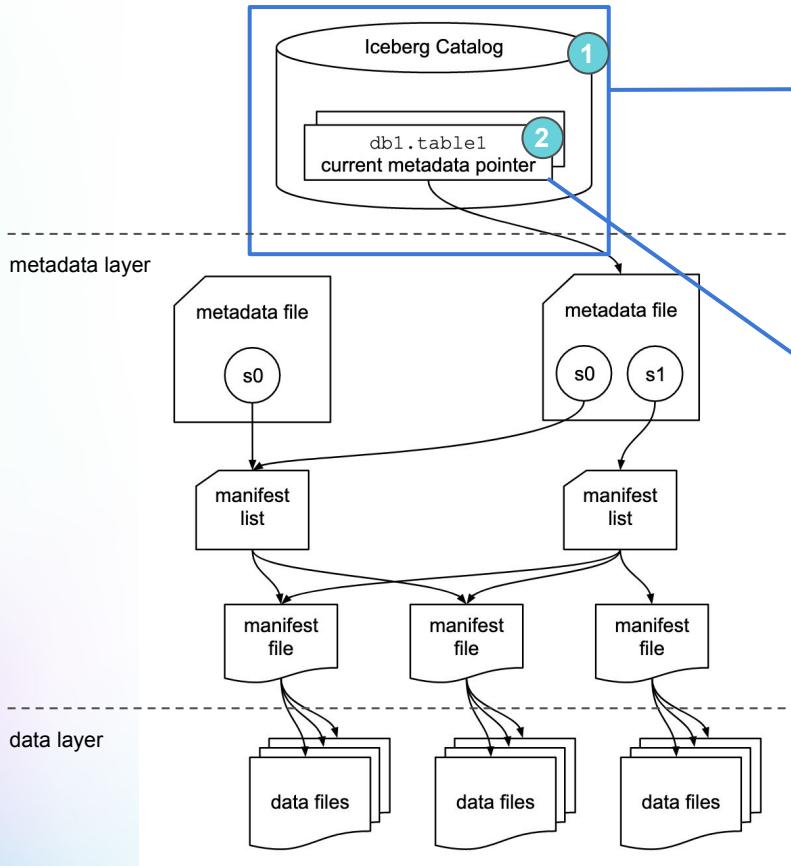


data layer

- **Overview of the components**
- **Summary of the read path (SELECT)**

#

# Iceberg components: Catalog



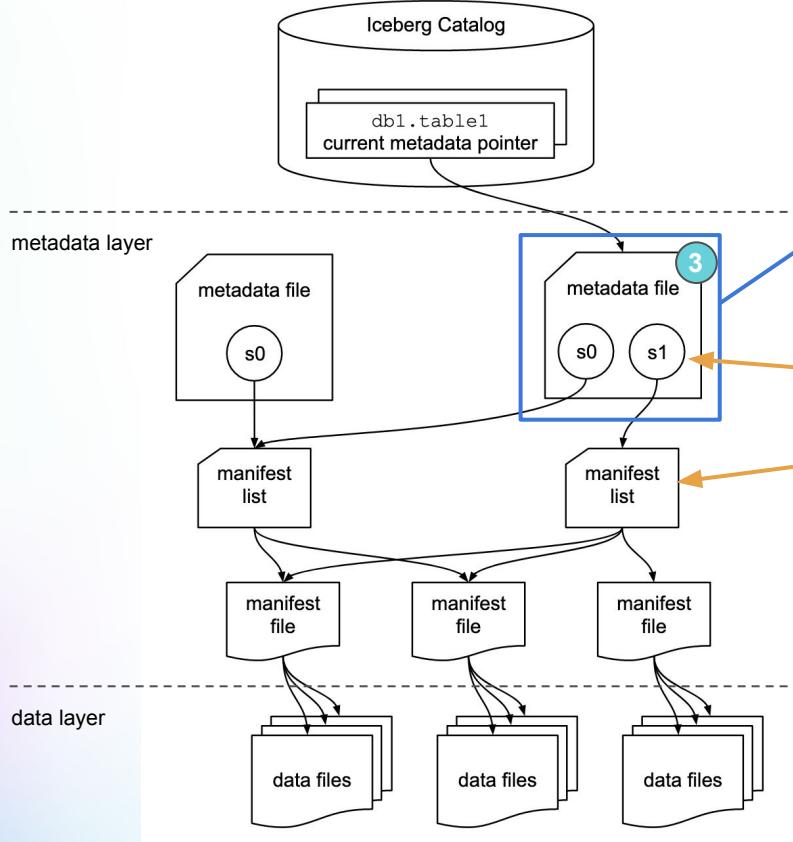
## Iceberg Catalog

- A store that houses the current metadata pointer for Iceberg tables
- Must support atomic operations for updating the current metadata pointer (e.g. HDFS, HMS, Nessie)

## table1's current metadata pointer

- Mapping of table name to the location of current metadata file

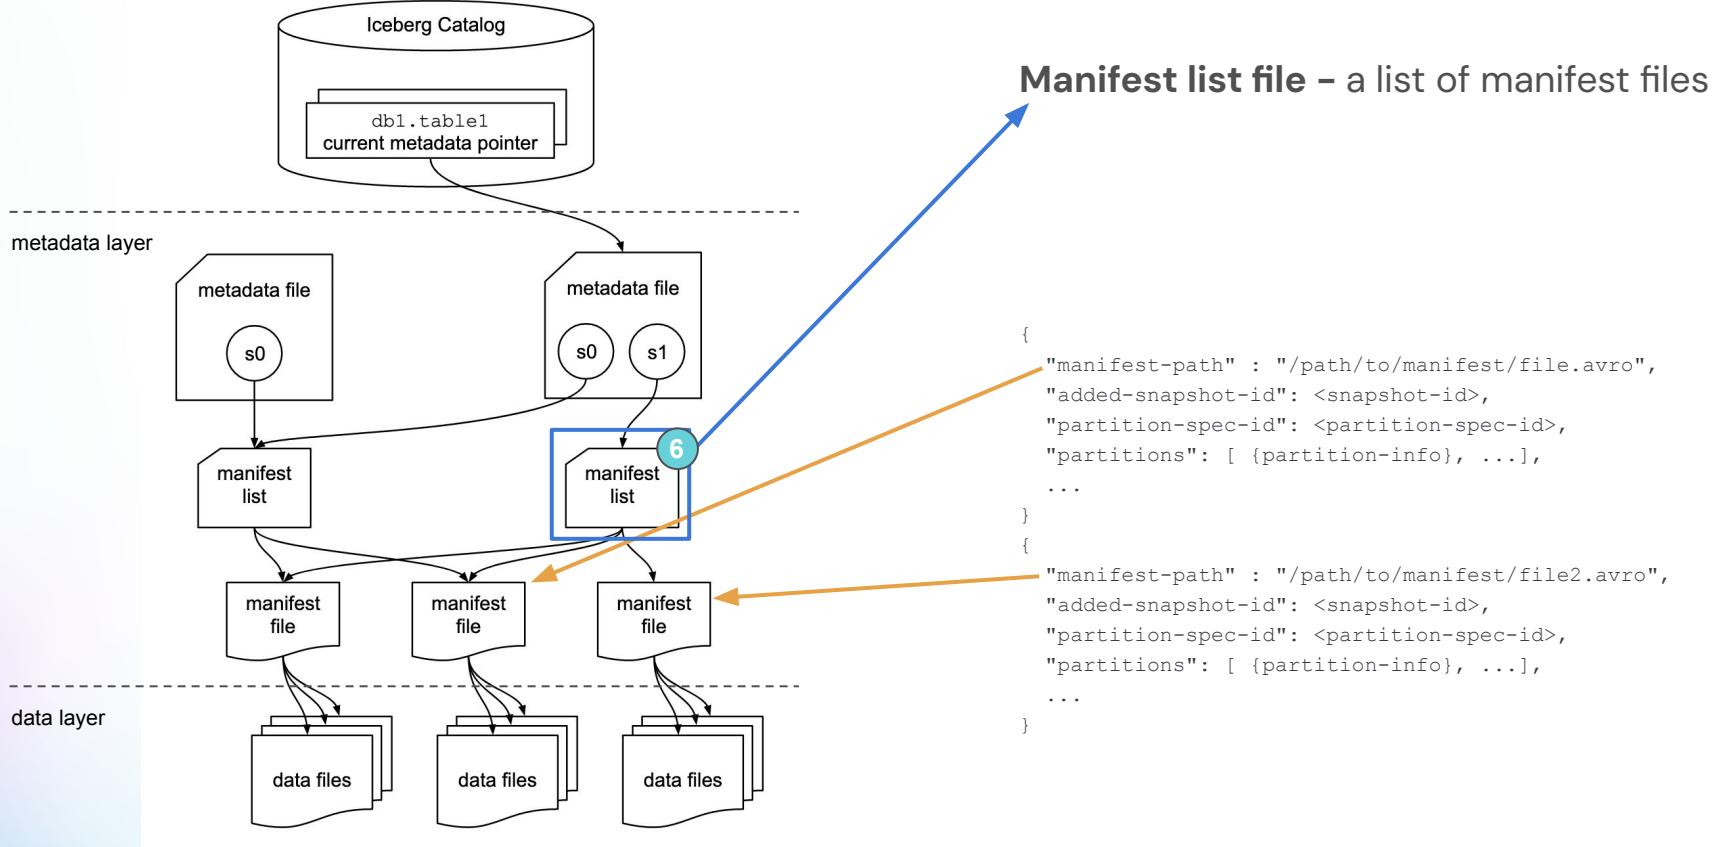
# Iceberg components: Metadata File



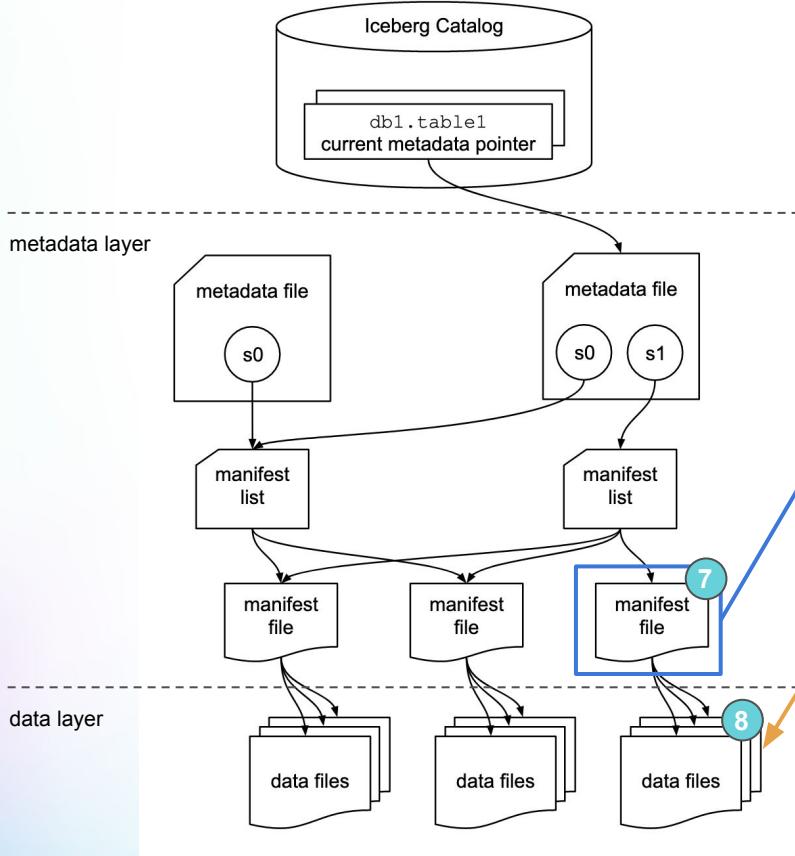
**Metadata file** – stores metadata about a table at a certain point in time

```
{  
    "table-uuid" : "<uuid>",  
    "location" : "/path/to/table/dir",  
    "schema": {...},  
    "partition-spec": [ {<partition-details>} , ... ],  
    4 "current-snapshot-id": <snapshot-id>,  
    "snapshots": [ {  
        "snapshot-id": <snapshot-id>  
        "manifest-list": "/path/to/manifest/list.avro"  
    } , ... ],  
    ...  
}
```

# Iceberg components: Manifest List



# Iceberg components: Manifest file



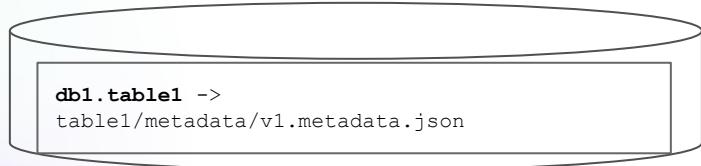
**Manifest file** – a list of data files, along with details and stats about each data file

```
{  
  "data-file": {  
    "file-path": "/path/to/data/file.parquet",  
    "file-format": "PARQUET",  
    "partition": {"<part-field>": {"<data-type>": <value>}},  
    "record-count": <num-records>,  
    "null-value-counts": [{  
      "column-index": "1", "value": 4  
    }, ...],  
    "lower-bounds": [{  
      "column-index": "1", "value": "aaa"  
    }, ...],  
    "upper-bounds": [{  
      "column-index": "1", "value": "eee"  
    }, ...]  
  }  
  ...  
}  
{  
  ...  
}
```

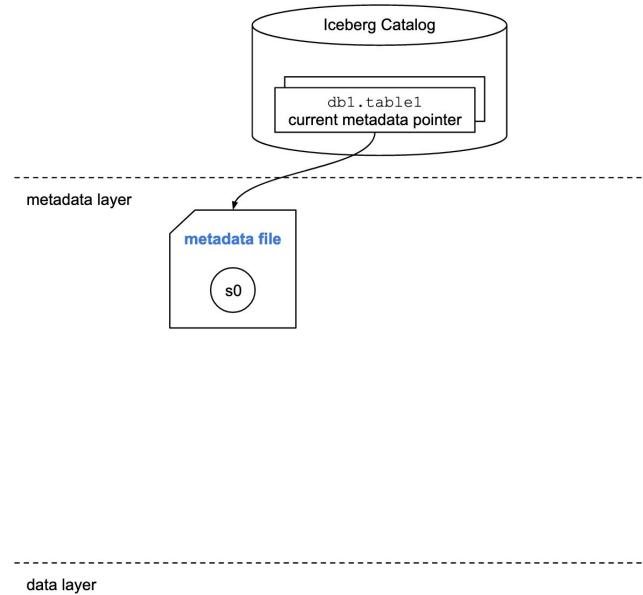
# Let's Look Under the Covers

# CREATE TABLE

```
CREATE TABLE db1.table1 (
    order_id bigint,
    customer_id bigint,
    order_amount DECIMAL(10, 2),
    order_ts TIMESTAMP
)
USING iceberg
PARTITIONED BY ( hour(order_ts) );
```

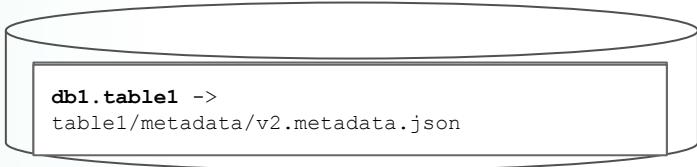


```
table1/
|- metadata/
|   |- v1.metadata.json
|- data/
```

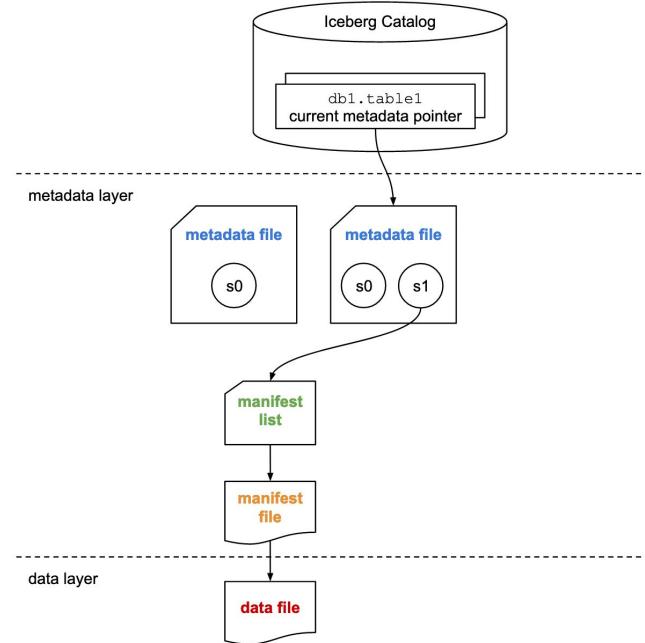


# INSERT

```
INSERT INTO db1.table1 VALUES (  
    123,  
    456,  
    36.17,  
    '2021-01-26 08:10:23'  
) ;
```

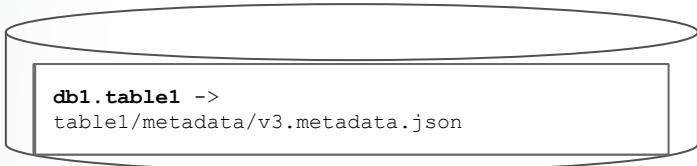


```
table1/  
| - metadata/  
|   | - v1.metadata.json  
|   | - v2.metadata.json  
|   | - snap-2938-1-4103.avro  
|   | - d8f9-ad19-4e.avro  
| - data/  
|   | - order_ts_hour=2021-01-26-08/  
|       | - 00000-5-cae2d.parquet
```

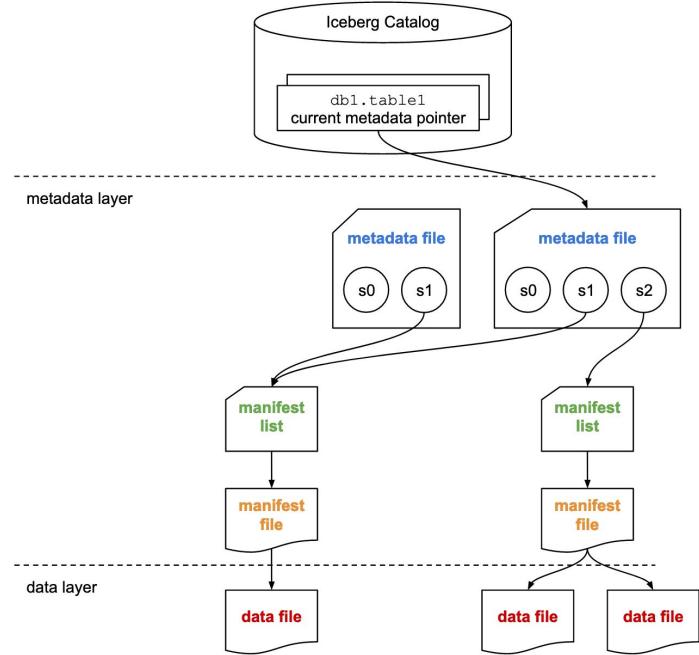


# UPSERT

```
MERGE INTO db1.table1
USING ( SELECT * FROM table1_stage ) s
ON table1.order_id = s.order_id
WHEN MATCHED THEN UPDATE table1.order_amount = s.order_amount
WHEN NOT MATCHED THEN INSERT *
```

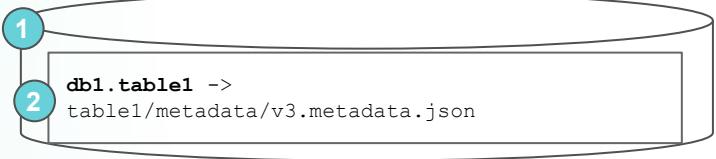


```
table1/
|- metadata/
|   |- v1.metadata.json
|   |- v2.metadata.json
|   |- v3.metadata.json
|   |- snap-29c8-1-b103.avro
|   |- snap-9fa1-3-16c3.avro
|   |- d8f9-ad19-4e.avro
|   |- 0d9a-98fa-77.avro
|- data/
  |- order_ts_hour=2021-01-26-08/
    |- 00000-5-cae2d.parquet
    |- 00000-1-aef71.parquet
  |- order_ts_hour=2021-01-27-10/
    |- 00000-3-0fa3a.parquet
```

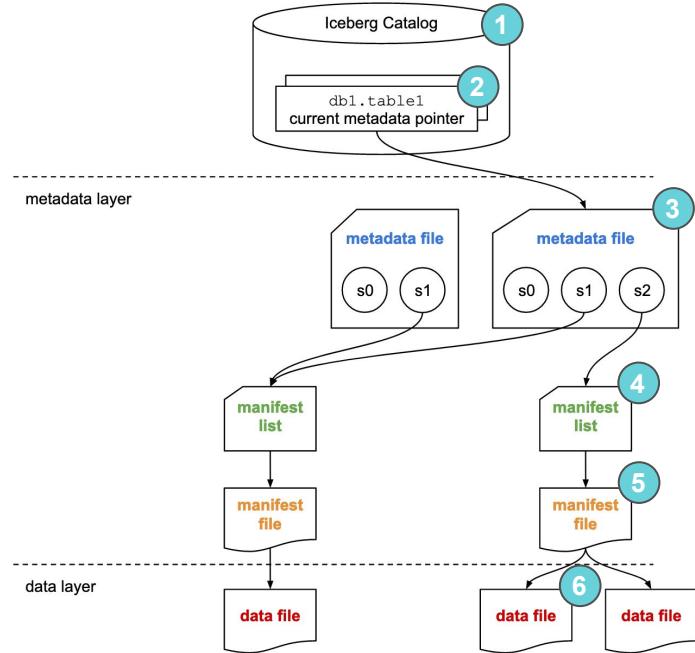


# READ

```
SELECT *\nFROM db1.table1\nWHERE order_ts = DATE '2021-01-26'
```

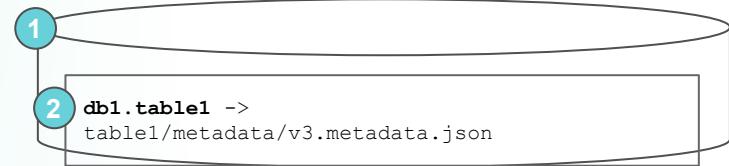


```
table1/\n|- metadata/\n|  |- v1.metadata.json\n|  |- v2.metadata.json\n|  3 v3.metadata.json\n|  |- snap-29c8-1-b103.avro\n|  4 snap-9fa1-3-16c3.avro\n|  |- d8f9-ad19-4e.avro\n|  5 0d9a-98fa-77.avro\n|- data/\n  |- order_ts_hour=2021-01-26-08/\n    |- 00000-5-cae2d.parquet\n    |  6 00000-1-aef71.parquet\n    |- order_ts_hour=2021-01-27-10/\n      |- 00000-3-0fa3a.parquet
```

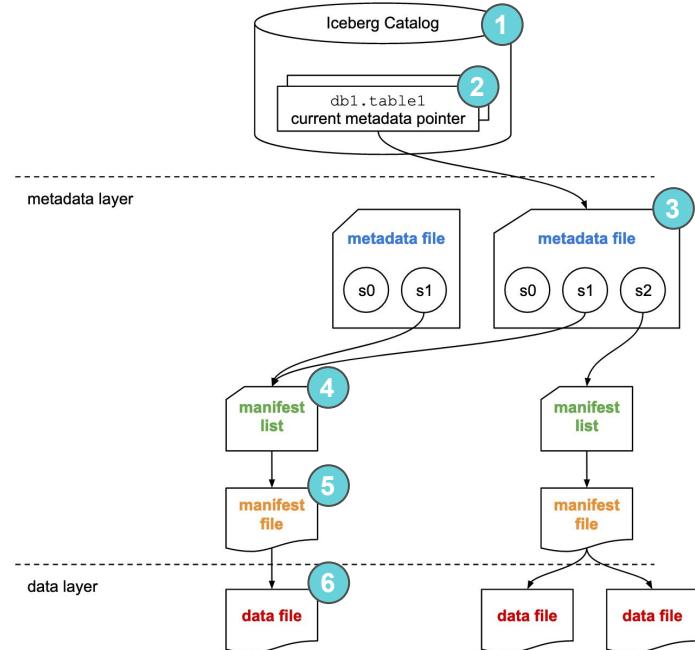


# TIME TRAVEL

```
SELECT *\nFROM db1.table1 AS OF '2021-05-26 09:30:00'\n-- (timestamp is from before MERGE INTO operation)
```



```
table1/\n  |- metadata/\n  |  |- v1.metadata.json\n  |  |- v2.metadata.json\n  |  3 v3.metadata.json\n  |  |- snap-29c8-1-b103.avro\n  |  |- snap-9fa1-3-16c3.avro\n  |  |- d8f9-ad19-4e.avro\n  |  5 0d9a-98fa-77.avro\n  |- data/\n    |- order_ts_hour=2021-01-26-08/\n      |  6 00000-5-cae2d.parquet\n      |  |- 00000-1-aef71.parquet\n      |- order_ts_hour=2021-01-27-10/\n        |  |- 00000-3-0fa3a.parquet
```



# Case Study: Atlas

- Historical Atlas data:
  - Time-series metrics from Netflix runtime systems
  - 1 month: 2.7 million files in 2,688 partitions
  - **Problem: cannot process more than a few days of data**
- Sample query:

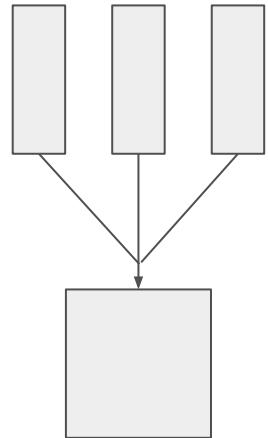
```
select distinct tags['type'] as type
from iceberg.atlas
where
    name = 'metric-name' and
    date > 20180222 and date <= 20180228
order by type;
```

# Case Study: Atlas Performance

- Hive table – with Parquet filters:
  - 400k+ splits per day, not combined
  - EXPLAIN query: **9.6 min (planning wall time)**
- Iceberg table – partition data filtering:
  - 15,218 splits, combined
  - **13 min** (wall time) / 61.5 hr (task time) / **10 sec** (planning)
- Iceberg table – partition and min/max filtering:
  - 412 splits
  - **42 sec** (wall time) / 22 min (task time) / 25 sec (planning)

# Enabled by this design: Compaction

- Asynchronously compact small files into fewer larger files
- It being asynchronous helps balance the write-side and read-side trade-offs
- Input and output of compaction jobs can be different file types
  - E.g. avro from streaming writes, compacted into larger parquet files for analytics
- Scheduling/triggering and the actual compaction work is done by external tools
  - Scheduling/triggering: scheduler, workflow tool, etc.
  - Compaction work execution: processing engine (e.g. Spark, Dremio)



# Design Benefits



- **Efficiently make smaller updates**
  - Make changes at the file level
- **Snapshot isolation for transactions**
  - Reads and writes don't interfere with each other and all writes are atomic
  - Concurrent writes
- **Faster planning and execution**
  - List of files defined on the write-side
  - Column stats in manifest files used to eliminate files
- **Reliable metrics for CBOs (vs hive..)**
  - Done on write instead of "infrequent" expensive read job
- **Abstract the physical, expose a logical view**
  - Hidden partitioning
  - Compaction
  - Tables can change over time
  - DE can transparently experiment with table layout
- **Rich schema evolution support**
- **All engines see changes immediately**

# Additional Resources

- [iceberg.apache.org](http://iceberg.apache.org)
- [iceberg.apache.org/blogs/](http://iceberg.apache.org/blogs/)
- [dremio.com/subsurface/apache-iceberg/](http://dremio.com/subsurface/apache-iceberg/)
- Get hands on – [iceberg.apache.org/getting-started/](http://iceberg.apache.org/getting-started/)

# Thank You



Presenter: Dipankar



@dipankartnt