


CS 417 – DISTRIBUTED SYSTEMS

# Week 11: Content Delivery

## Part 3: Memory Caching

Paul Krzyzanowski



© 2022 Paul Krzyzanowski. No part of this content may be reproduced or reposted in whole or in part in any manner without the permission of the copyright owner.

# Caching

## Purpose of a cache

- Temporary storage to increase data access speeds
- Increase effective bandwidth by caching most frequently used data

## Store raw data from slow devices

- Memory cache on CPUs
- Buffer cache in operating system
- Chubby file data and metadata
- GFS master caches all metadata in memory

## Store computed data

- Avoid the need to look the same thing up again
  - Results of database queries or file searches
  - Spark RDDs in memory

# What would you use a caching service for?

## Cache user session state on web application servers

No need to have user come back to the same computer

## Cache user preferences, shopping carts, etc.

Avoid repeated database lookups (e.g., key-value data)

## Cache rendered web pages

Avoid re-processing server-side includes, JSP/ASP/PHP code

## Cache precomputed results

Avoid re-computing data that gets reused  
(Spark RDDs, news posts, inventory status, ...)

# Distributed In-Memory Caching as a Service

## A network memory-based caching service

Shared by many – typically used by front-end services

## Stores frequently-used (key, value) data

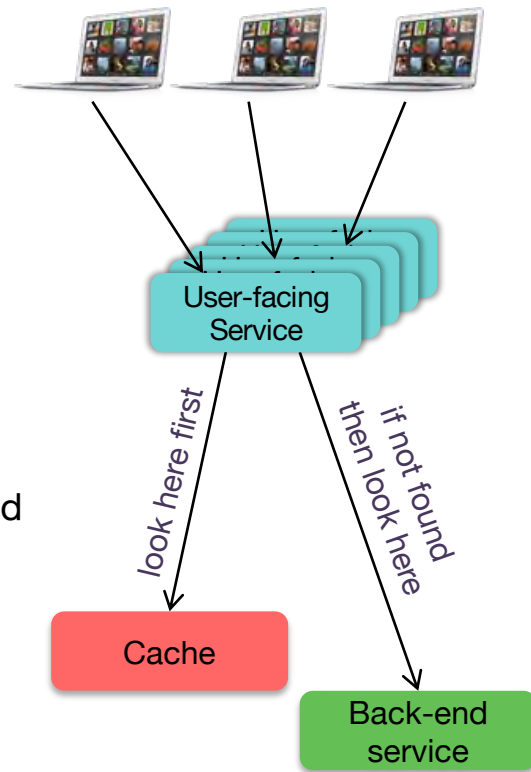
Old data gets evicted

## General purpose

Not tied to a specific back-end service

## Not transparent (usually)


Because it's a general-purpose service, the programmer gets involved



# Deployment Models


## Separate caching server

One or more computers whose sole purpose is to provide a caching service



## Or share cache memory among servers

Take advantage of free memory from lightly-loaded nodes



# Example: memcached

Free & open source distributed memory caching

## Used by

Dropbox, Facebook, Wikipedia, Flickr, Twitter,  
YouTube, Instagram, Digg, Bebo, WordPress, Craigslist, ...

## Protocol

Binary & ASCII versions

## Client APIs for

Command line, C/C++, C#, Go, PHP, Java, Python, Ruby, Perl, Erlang, Lua, LISP,  
Windows/.NET, mySQL, PostgreSQL, ColdFusion, ...



[memcached.org](http://memcached.org)

# Memcached structure

## Key-Value store

- Cache is made up of { **key**, **value**, **expiration time**, **flags** }
- All access is  $O(1)$

## Client software

- Provided with a list of *memcached* servers
- Hash(key) chooses a server based on the *key*

## Server software

- Stores keys and values in an in-memory hash table
- Throw out old data when necessary
  - LRU cache and time-based expiration
  - Objects expire after a minute to ensure stale data is not returned
- **Servers are unaware of each other**

# Memcached API

## Commands sent over TCP (UDP also available)

Connection may be kept open indefinitely

### Commands

#### Storage

- Storage commands take an expiration time in seconds from current time or 0 = forever (but may be deleted)
- **set** – store data
- **add** – store data only if the server does not have data for the key
- **replace** – store data if the server does have data for the key
- **append** – add data after existing data
- **prepend** – add data before existing data
- **cas** – check & set: *store data only if no one else updated it since I fetched it* (cas = unique, 64-bit value associated with the item)

#### Retrieval

- **get** – retrieve one or more keys: returns *key, flags, bytes, and cas unique*



# Memcached API

## Commands

### Deletion

- `delete` *key*

### Increment/decrement

- Treat data as a 64-bit unsigned integer and add/subtract value
- `incr` *key value* – increment *key* by *value*
- `decr` *key value* – decrement *key* by *value*

### Update expiration

- `touch` *key exptime* – Update the expiration time

### Get Statistics

- `stats` – various options for reporting statistics

### Flush

- `flush_all` – clear the cache

# Redis

## Memory cache + in-memory database + message broker

Open source: see [redis.io](https://redis.io)

Text-based command interface

### Features

- Key-value store
- Transactions
- Publish/subscribe messaging
- Expiration of data
- Built-in replication
- Optional disk persistence
- Lua scripting (via EVAL command)
- Automatic partitioning with Redis Cluster



Used by: Twitter, GitHub, Weibo, Pinterest, Snapchat, Craigslist, Digg, StackOverflow, Flickr, Shopify, Hulu, Trello, Uber, Coinbase, ...

# Redis Data Types

## Strings

- Simplest type; only type supported in memcached)

## Lists

- Collections of strings sorted by order of insertion

## Sets

- Collections of unique, unsorted strings

## Sorted sets

- Every element is associated with a **score** (floating point number)
- Elements sorted by score
- Operations to retrieve ranges (e.g., top 10, bottom 10)

## Hashes

- Maps of fields associated with values (fields & values are strings)

## Bitmaps

- Commands to treat strings as bits (set/clear bits)

## HyperLogLogs

- Probabilistic data structure to estimate the cardinality of a set
  - Count # of unique items without storing the entire set of items
- Use a fixed amount of memory

# Redis as a memory cache: Timeouts & Evictions

## Set expiration for specific keys

- Associate a timeout with a key
- Key deleted after the timeout

```
SET mykey "hello"
```

```
EXPIRE mykey 10
```

*expire key in 10 seconds*

## Tell the cache to automatically evict (delete) old data

### Methods of eviction

- LRU (least recently used)
- LRU only for keys that have an expiration time
- Random
- Random only for keys that have an expiration time

# Redis as an in-memory database

## **MULTI**

- Mark the start of a transaction (operations queued until EXEC)

## **EXEC**

- Execute queued commands in a transaction

## **DISCARD**

- Abort transaction & revert to previous values

## **WATCH**

- Test-and-set behavior to ensure mutual exclusion
- Monitor keys to detect changes
- Abort if change takes place

# Redis as a message broker

## **Publish/subscribe model**

- Senders (publishers) do not send messages to specific receivers
- Messages go to channels
- Subscribers listen to one or more channels, receiving messages of interest

## **Allows for scalability and dynamic topology**

- Publishers do not know subscribers
- Subscribers do not know publishers

## **Support for pattern-based channels**

- Subscribe to all channel names matching a pattern

# Redis partitioning

Data can be partitioned across multiple computers

## Types of partitions

- Range partitioning
  - Use table that maps ranges to instances
- Hash partitioning
  - Based on hash(key): works with any key

## Who does the partitioning?

- Client-side partitioning
- Proxy-assisted partitioning
- Query forwarding by a Redis server

# The End