Netflix Play API

Why we built an Evolutionary Architecture

Suudhan Rangarajan (@suudhan)
Senior Software Engineer

NETFLIX

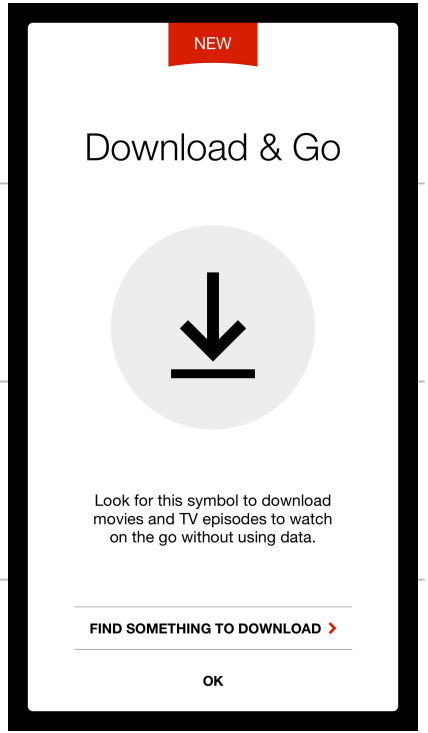

Q1 2016 Q2 2016 Q3 2016 Q4 2016

NETFLIX

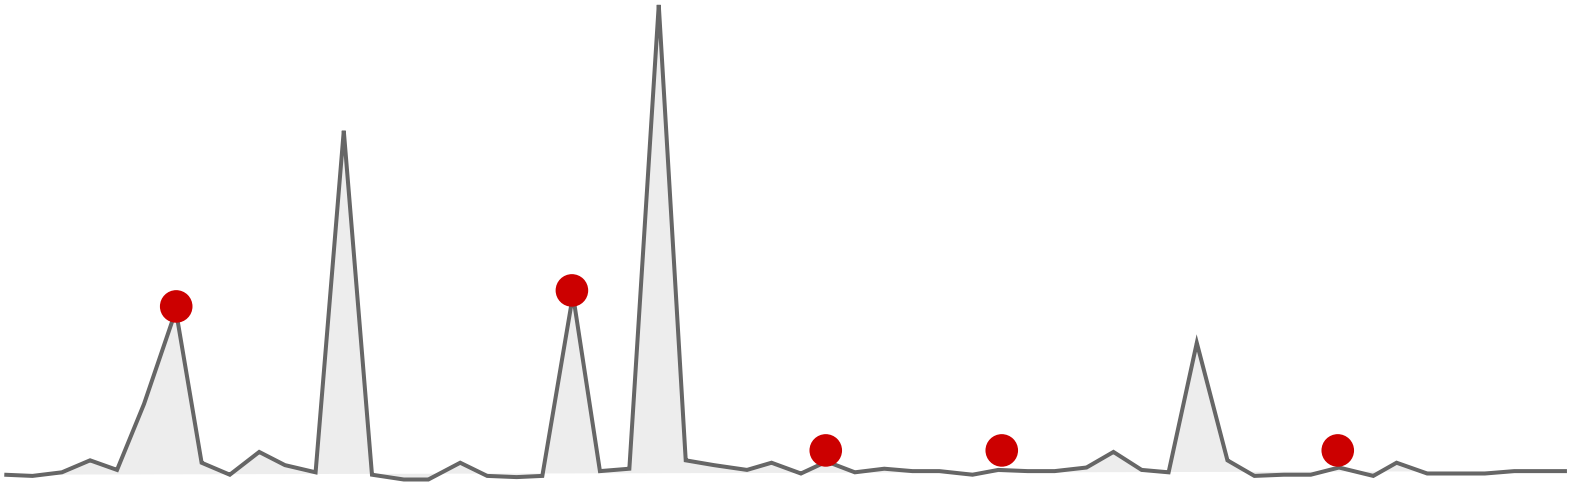

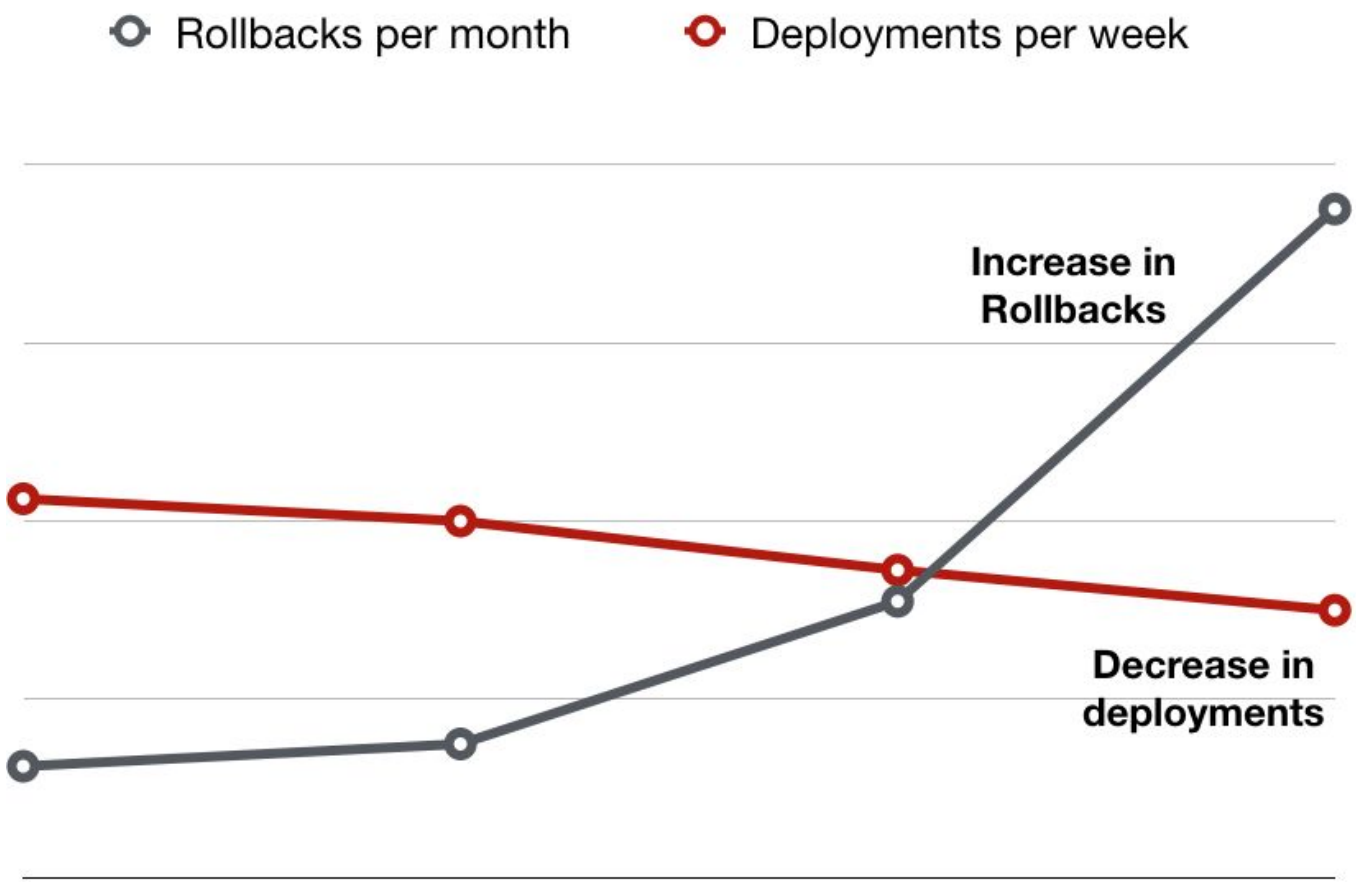
Rollbacks per month
Deployments per week
Increase in Rollbacks
Decrease in deployments

Netflix Play API

Why we built an Evolutionary Architecture

Suudhan Rangarajan (@suudhan)
Senior Software Engineer

NETFLIX

Previous Architecture Workflow

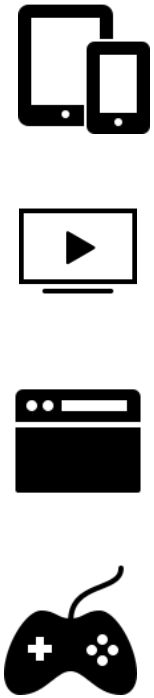

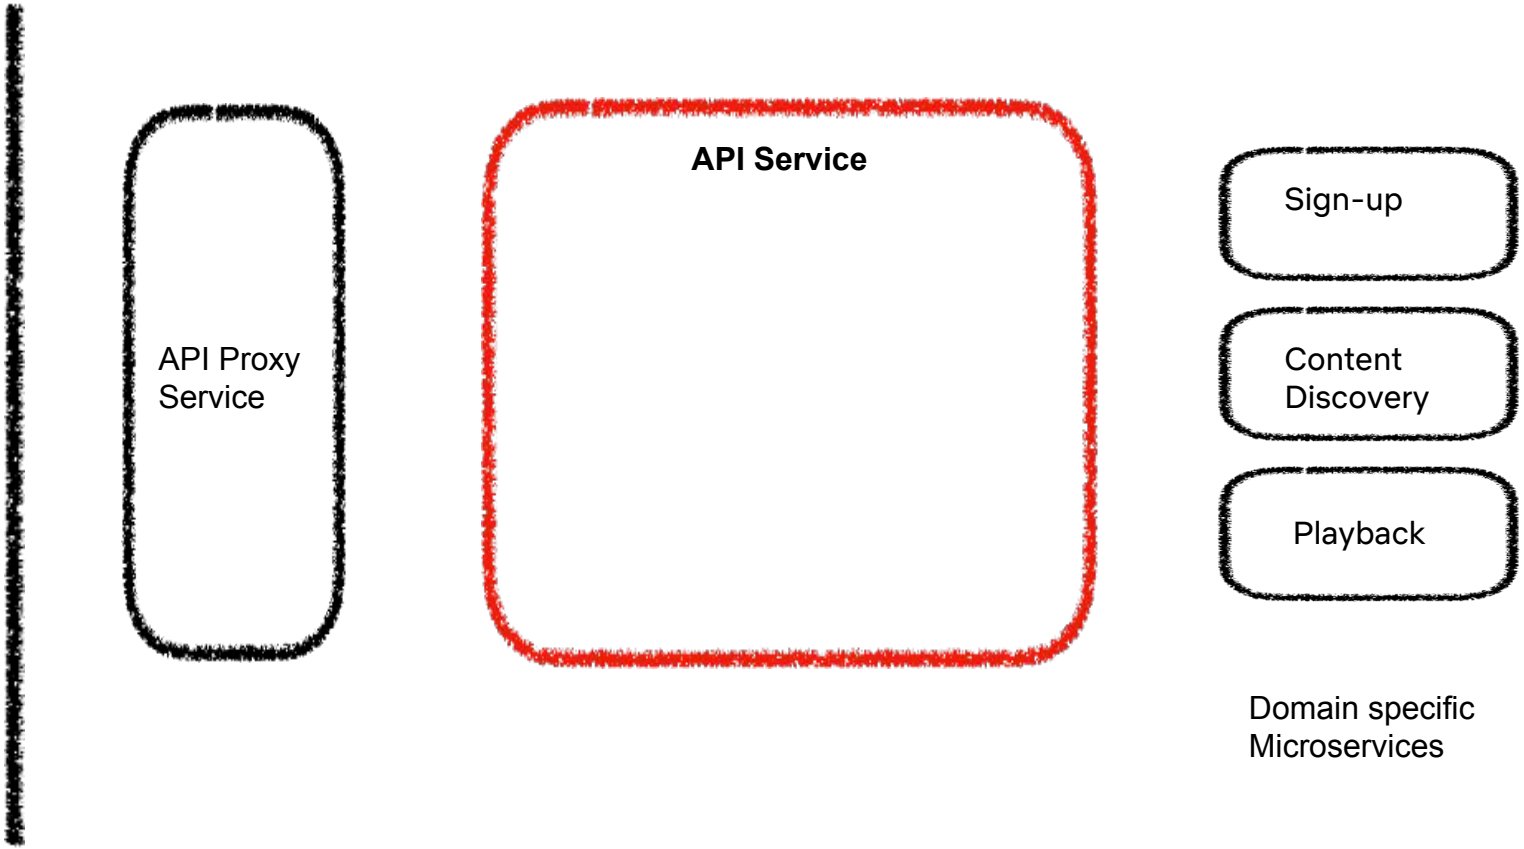

Devices

← Services hosted in AWS →

Signup Workflow

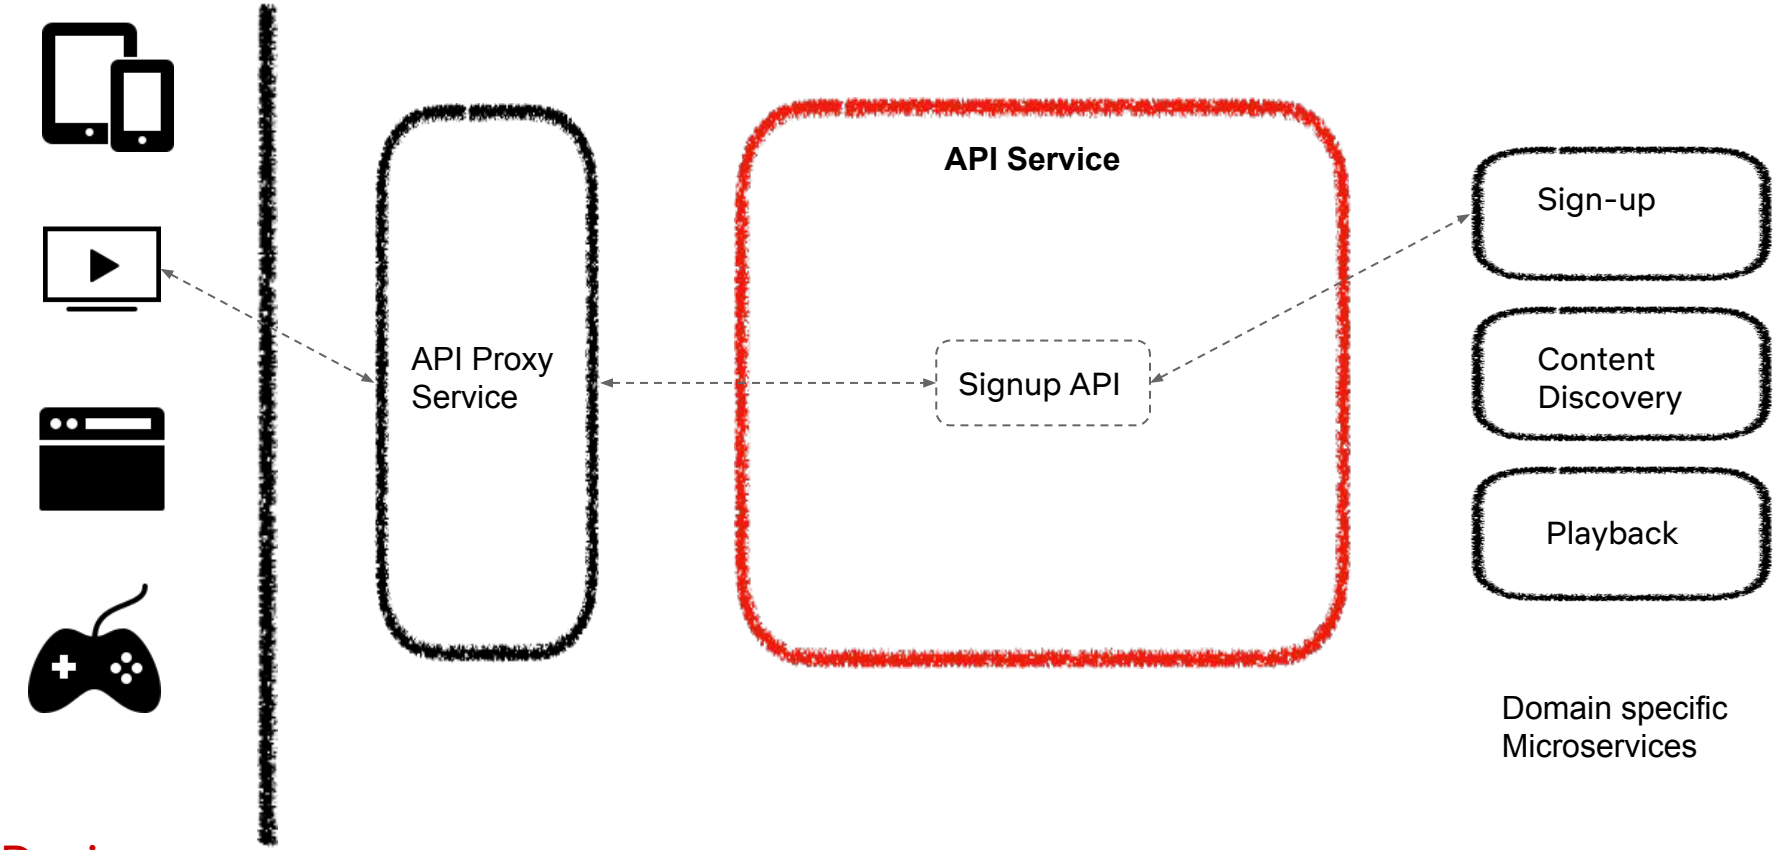

Content Discovery Workflow

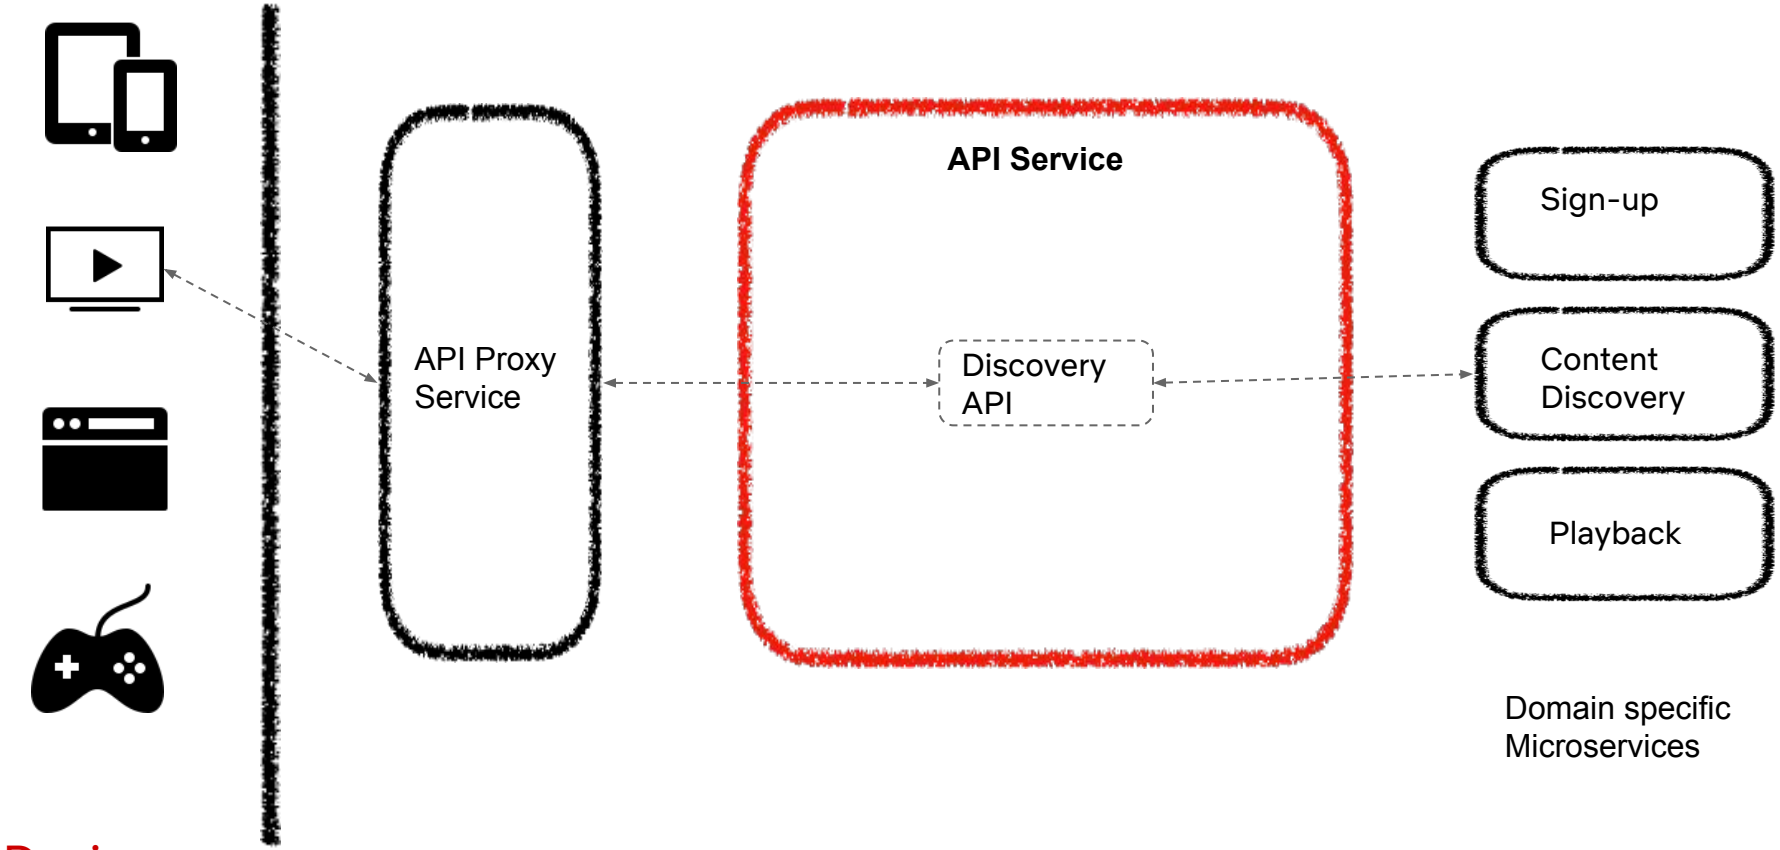

Playback Workflow

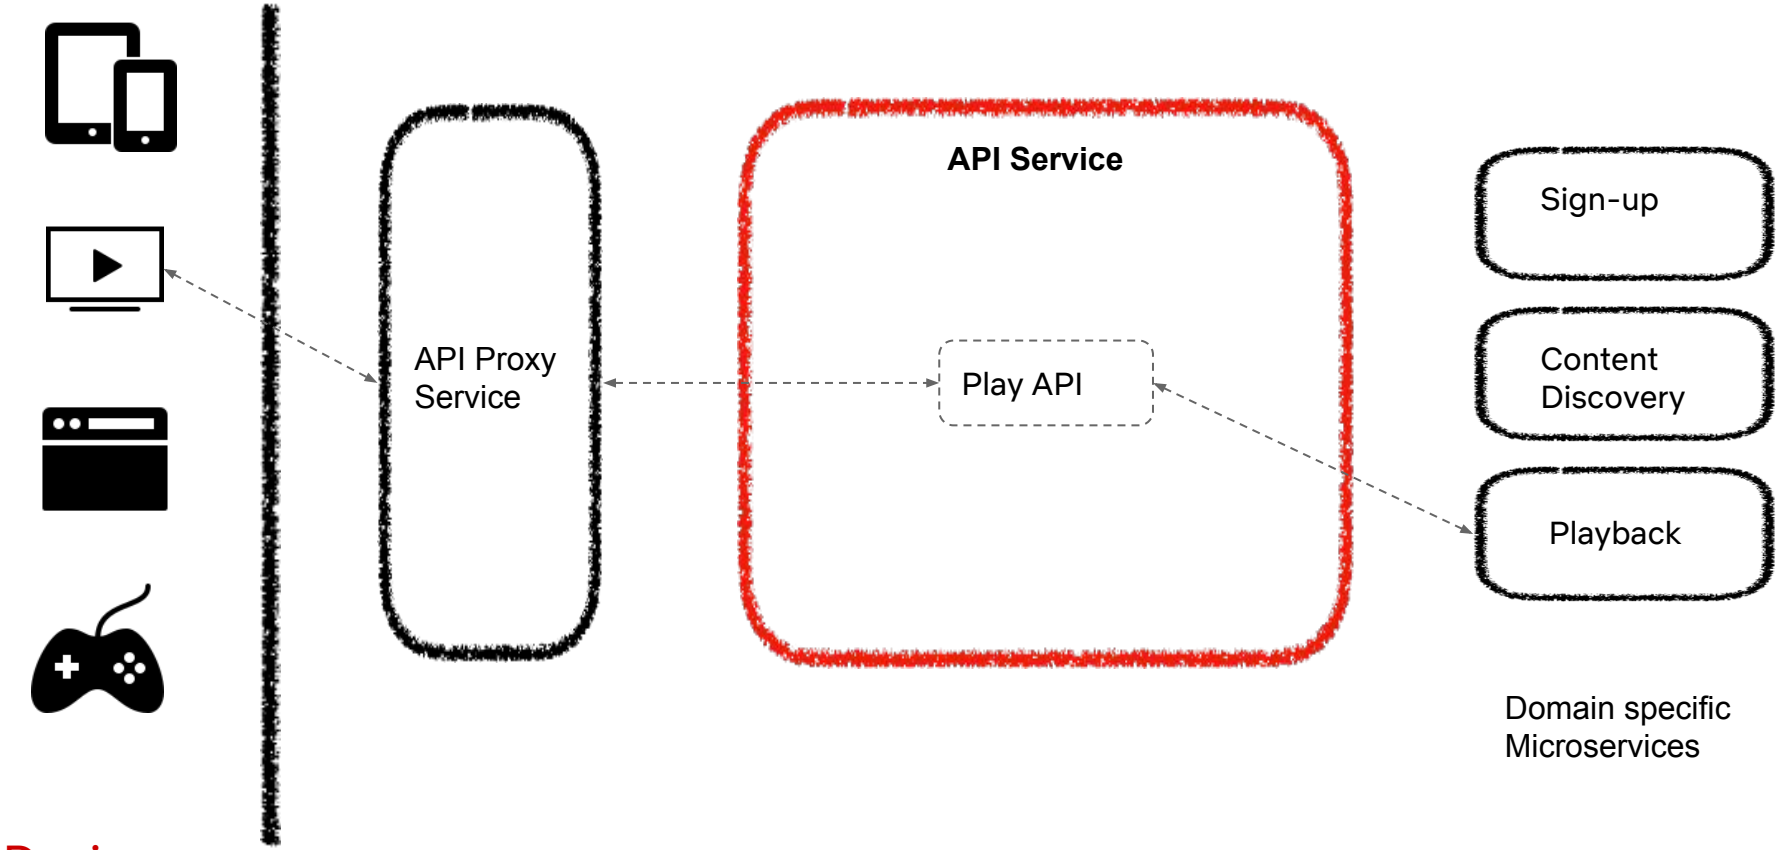

Previous Architecture

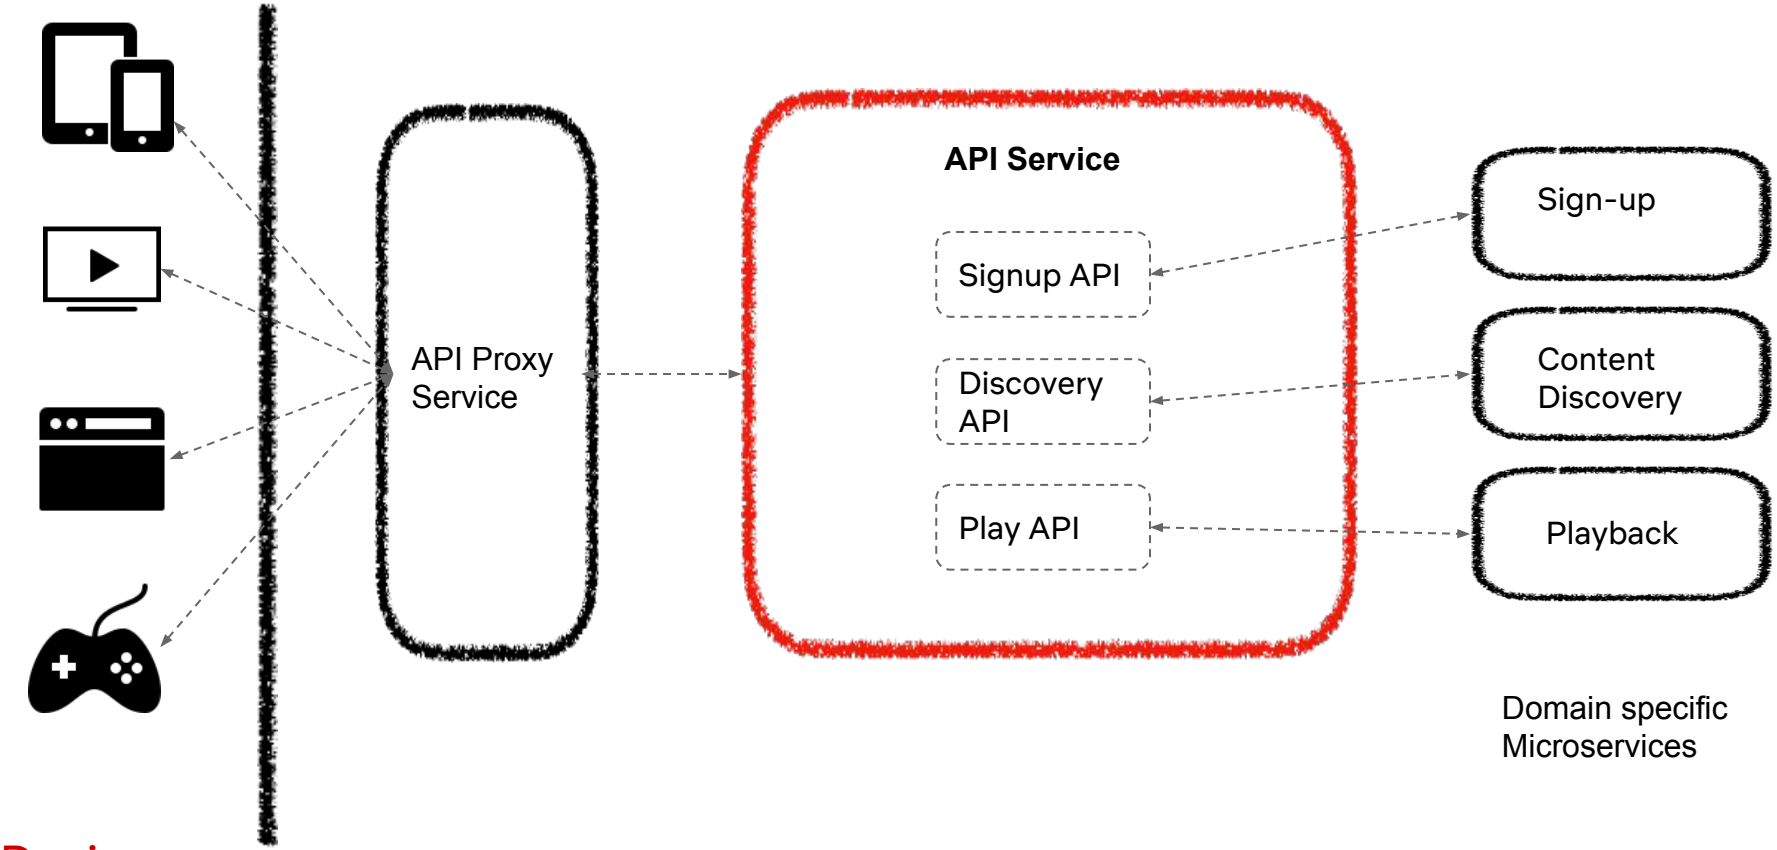

Identity

Type 1/2 Decisions

Evolvability

NETFLIX

Identity

Type 1/2 Decisions

Evolvability

NETFLIX

Start with WHY: Ask why your service exists

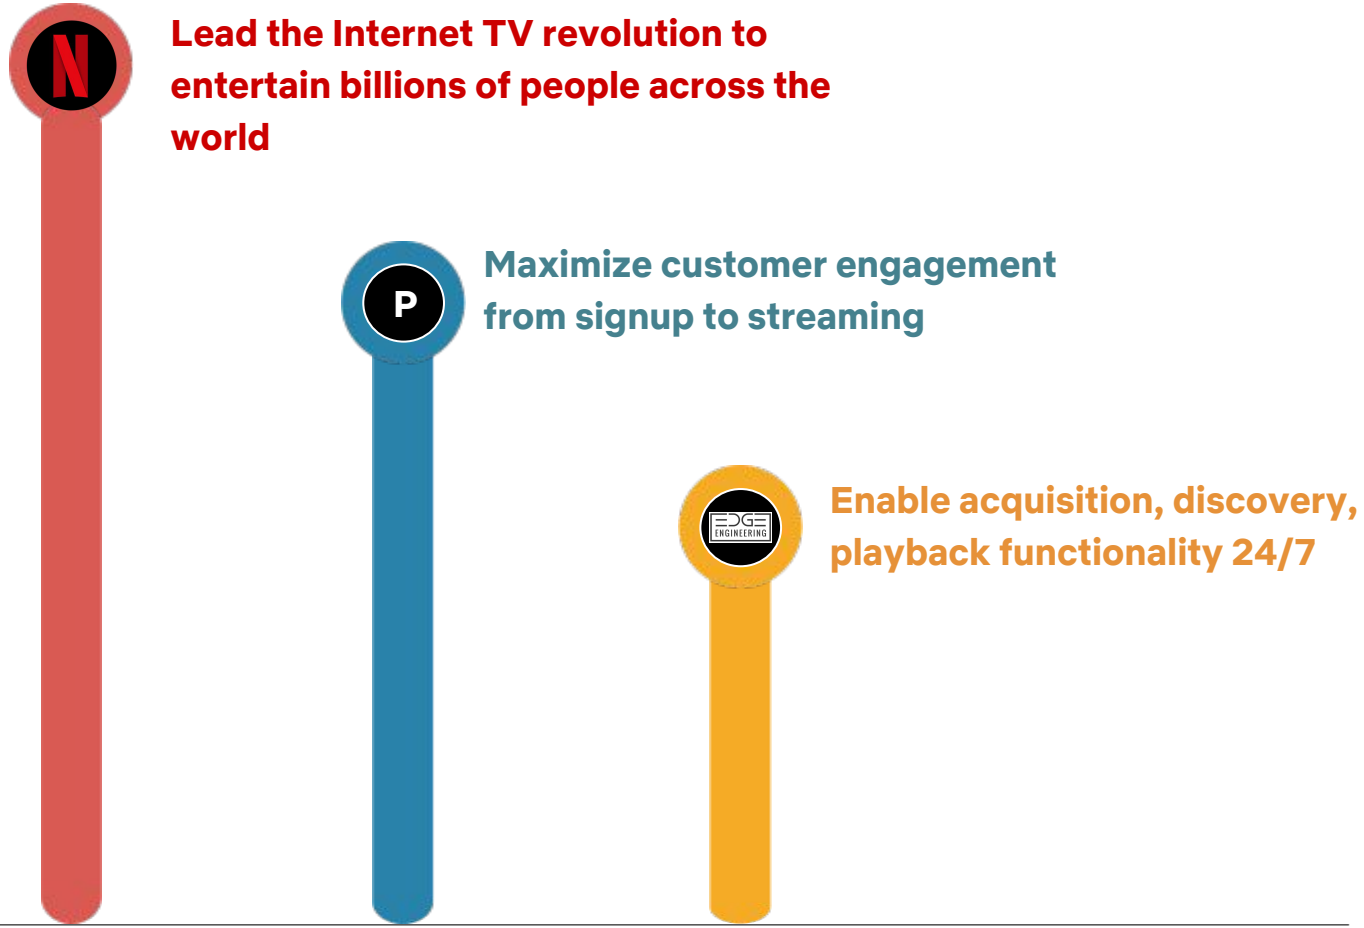
Lead the Internet TV revolution to entertain billions of people across the world
Maximize customer engagement from signup to streaming
Enable acquisition, discovery, playback functionality 24/7

NETFLIX

API Identity: Deliver Acquisition, Discovery and Playback functions with high availability

Single Responsibility Principle: Be wary of multiple-identities rolled up into a single service

Previous Architecture

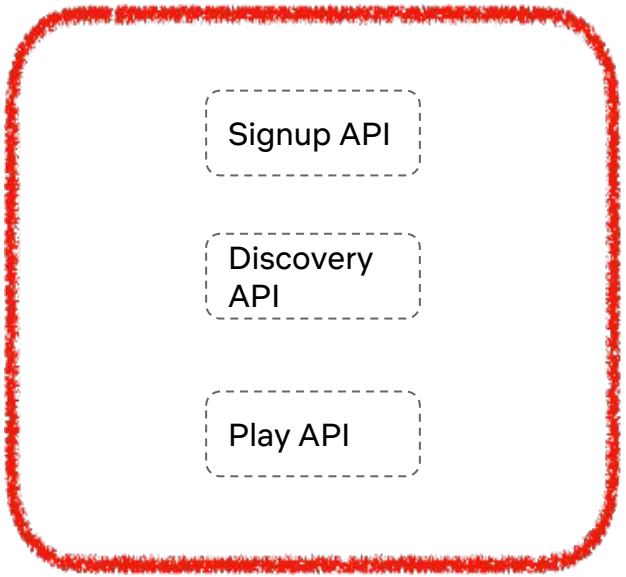

One API Service

Current Architecture

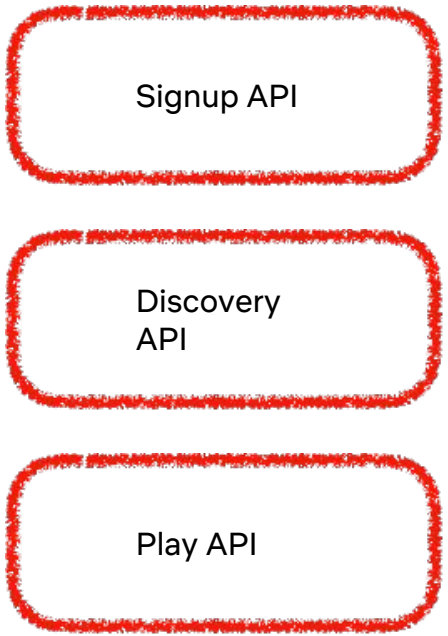

API Service Per function

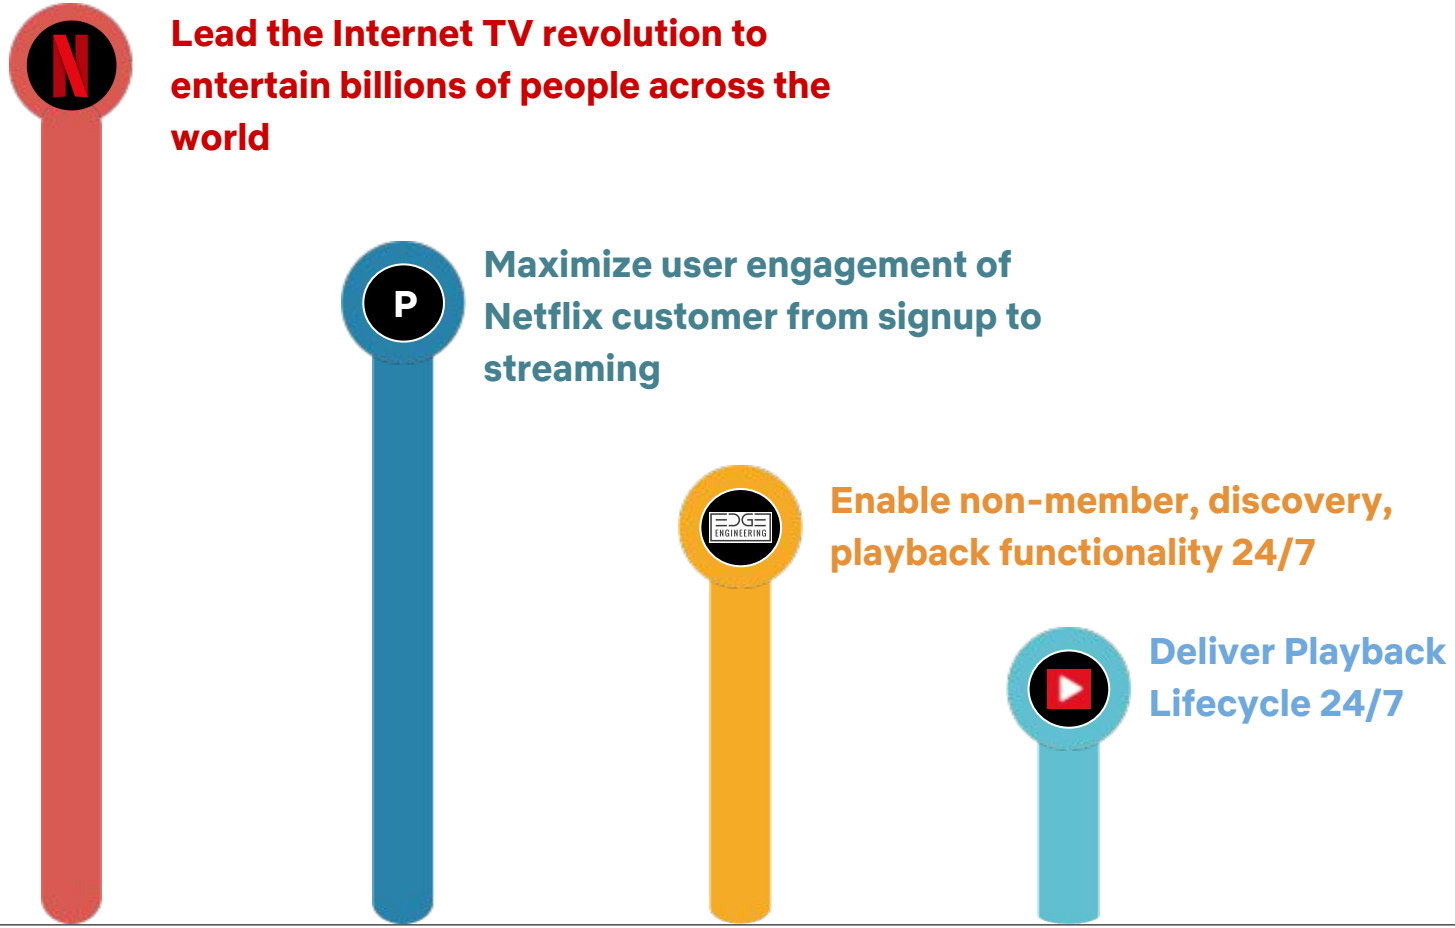
Lead the Internet TV revolution to entertain billions of people across the world
Maximize user engagement of Netflix customer from signup to streaming
Enable non-member, discovery, playback functionality 24/7
Deliver Playback Lifecycle 24/7

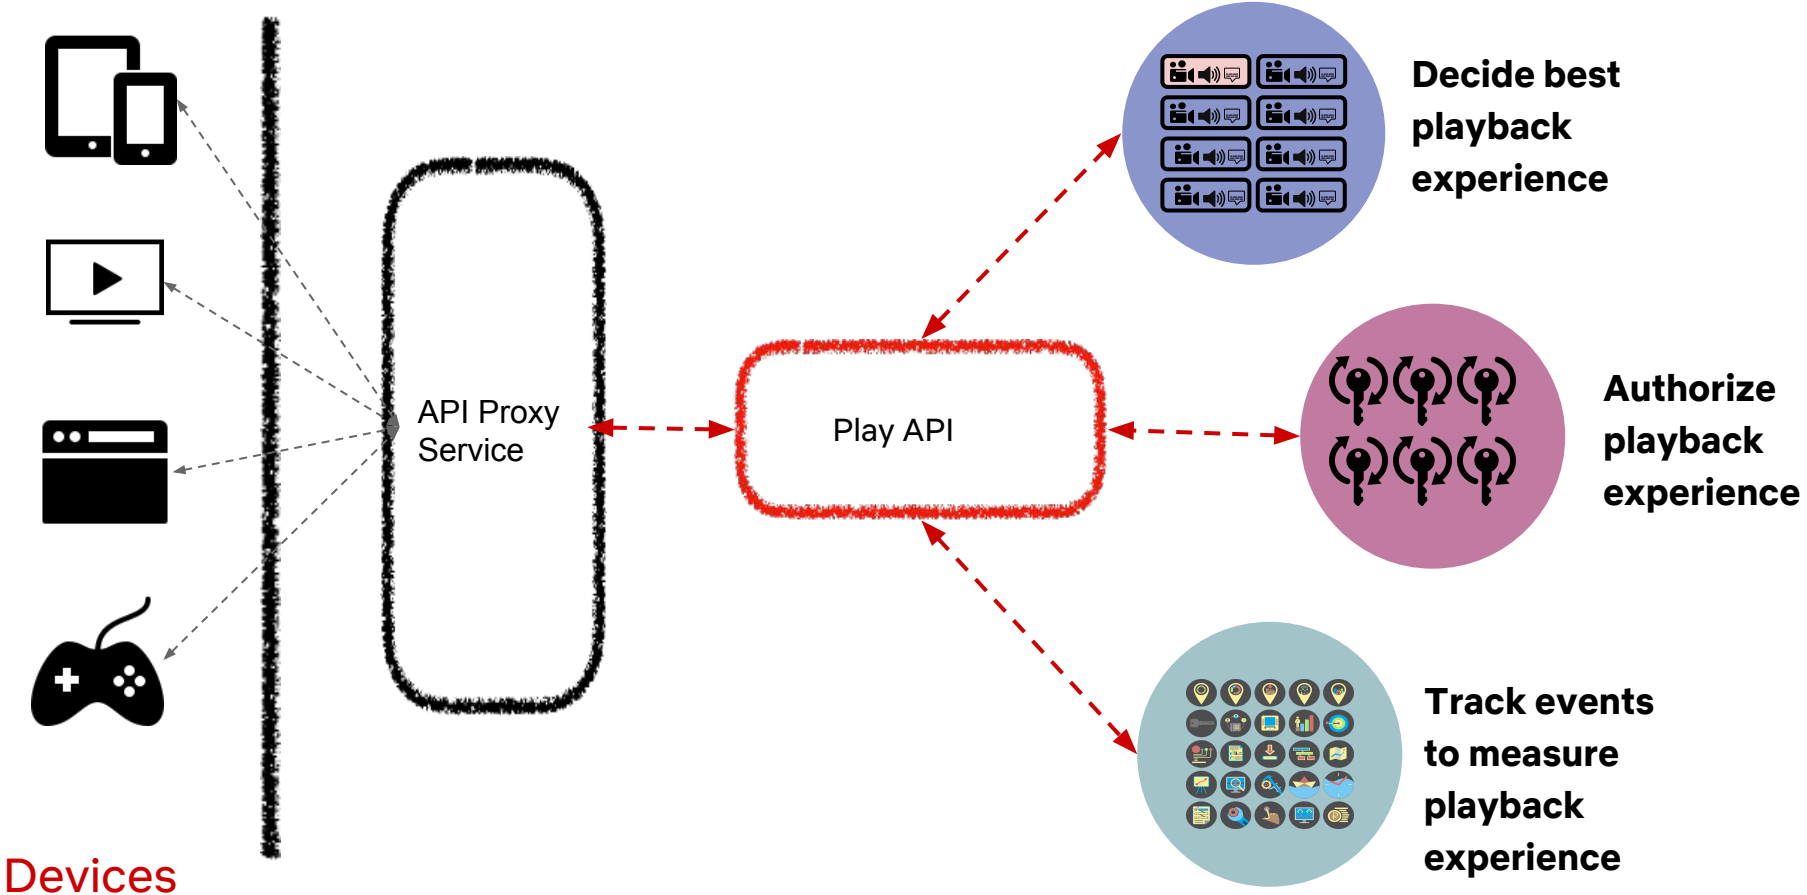
API Proxy Service
Play API
Decide best playback experience
Authorize playback experience
Track events to measure playback experience
Devices

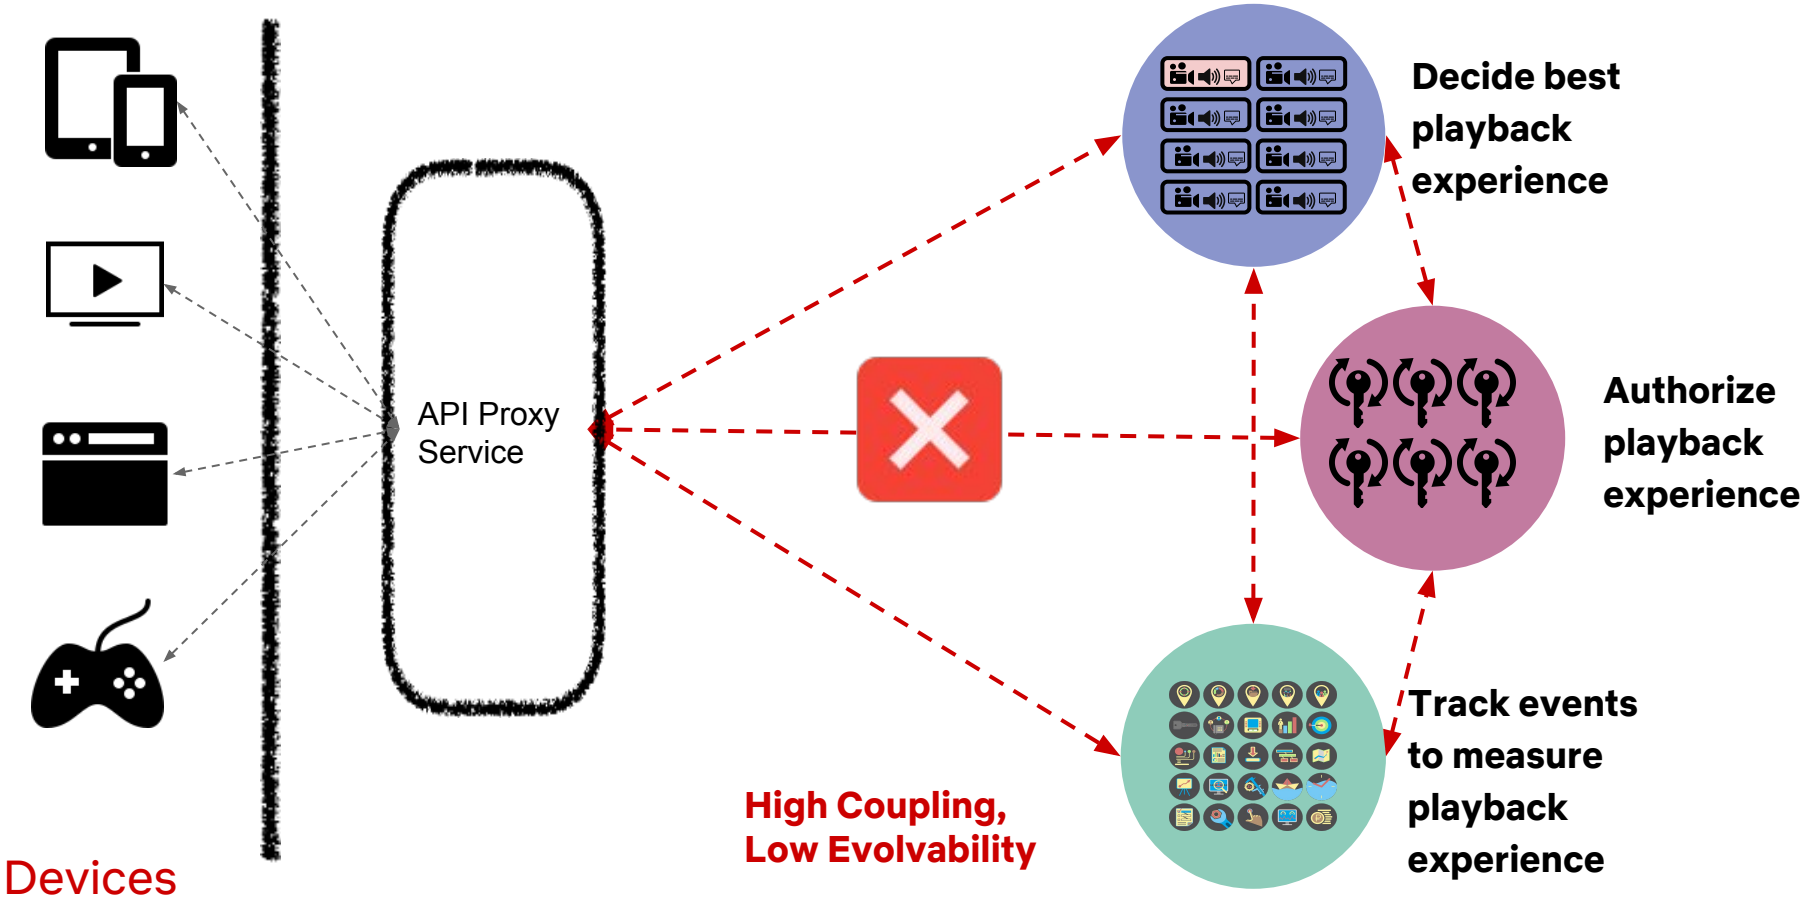
API Proxy
Service
Decide best
playback
experience
Authorize
playback
experience
Track events
to measure
playback
experience
High Coupling,
Low Evolvability
Devices

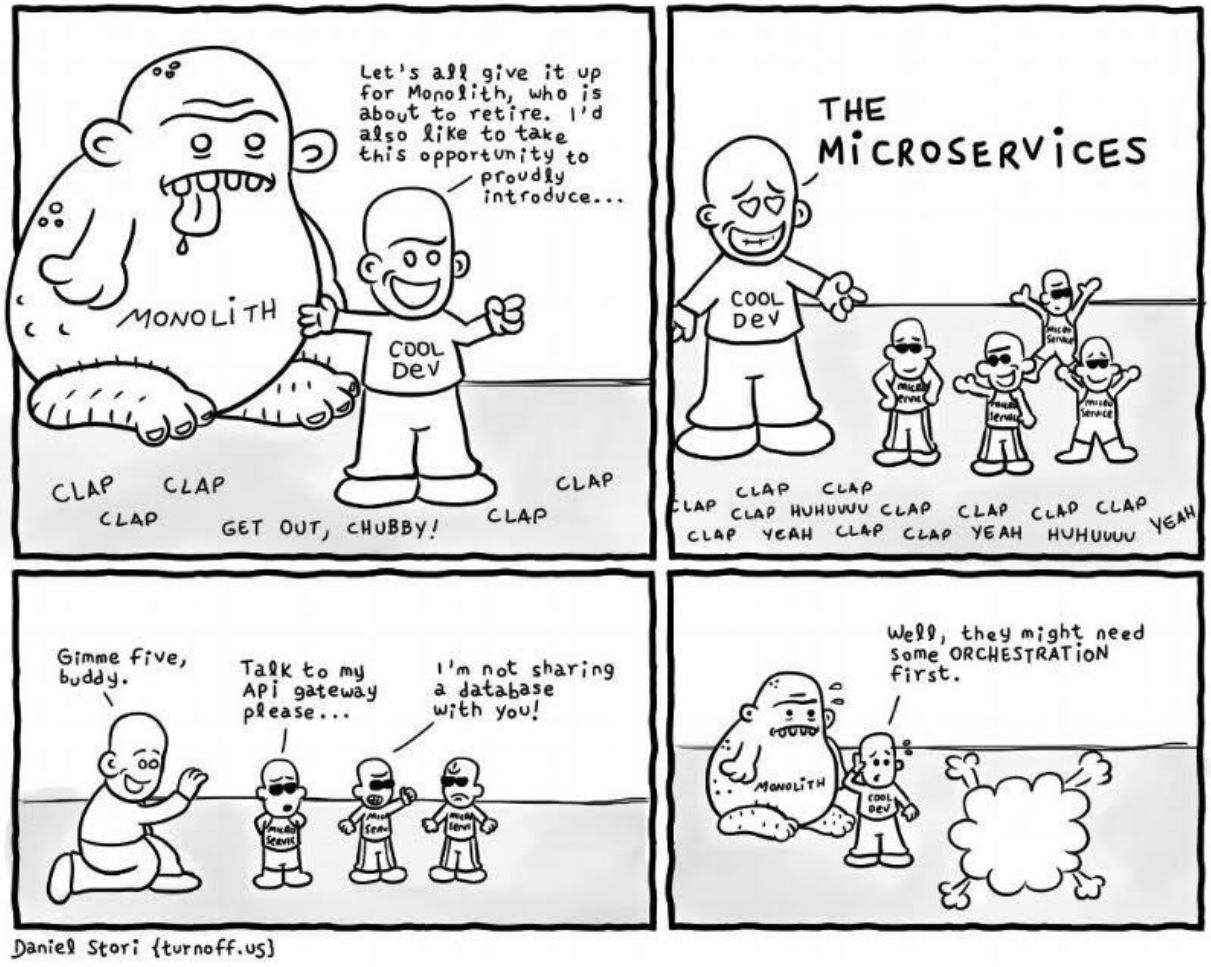

Daniel Stori {turnoff.us}

Play API Identity: Orchestrate Playback Lifecycle with stable abstractions

Guiding Principle: We believe in a simple singular identity for our services. The identity relates to and complements the identities of the company, organization, team and its peer services

Identity

Type 1/2 Decisions

Evolvability

NETFLIX

“Some decisions are consequential and irreversible or nearly irreversible – one-way doors – and these decisions must be made methodically, carefully, slowly, with great deliberation and consultation [...] We can call these Type 1 decisions...”

Quote from Jeff Bezos

“...But most decisions aren’t like that – they are changeable, reversible – they’re two-way doors. If you’ve made a suboptimal Type 2 decision, you don’t have to live with the consequences for that long [...] Type 2 decisions can and should be made quickly by high judgment individuals or small groups.”

Quote from Jeff Bezos

Three Type 1 Decisions to Consider

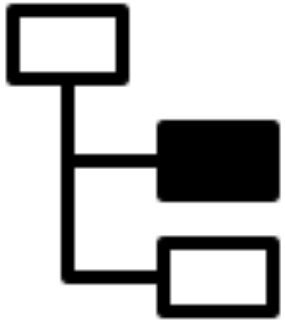

Appropriate Coupling

Synchronous & Asynchronous

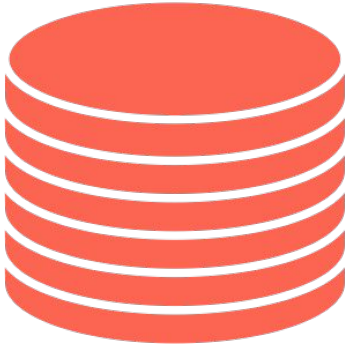

Data Architecture

Two types of Shared Libraries

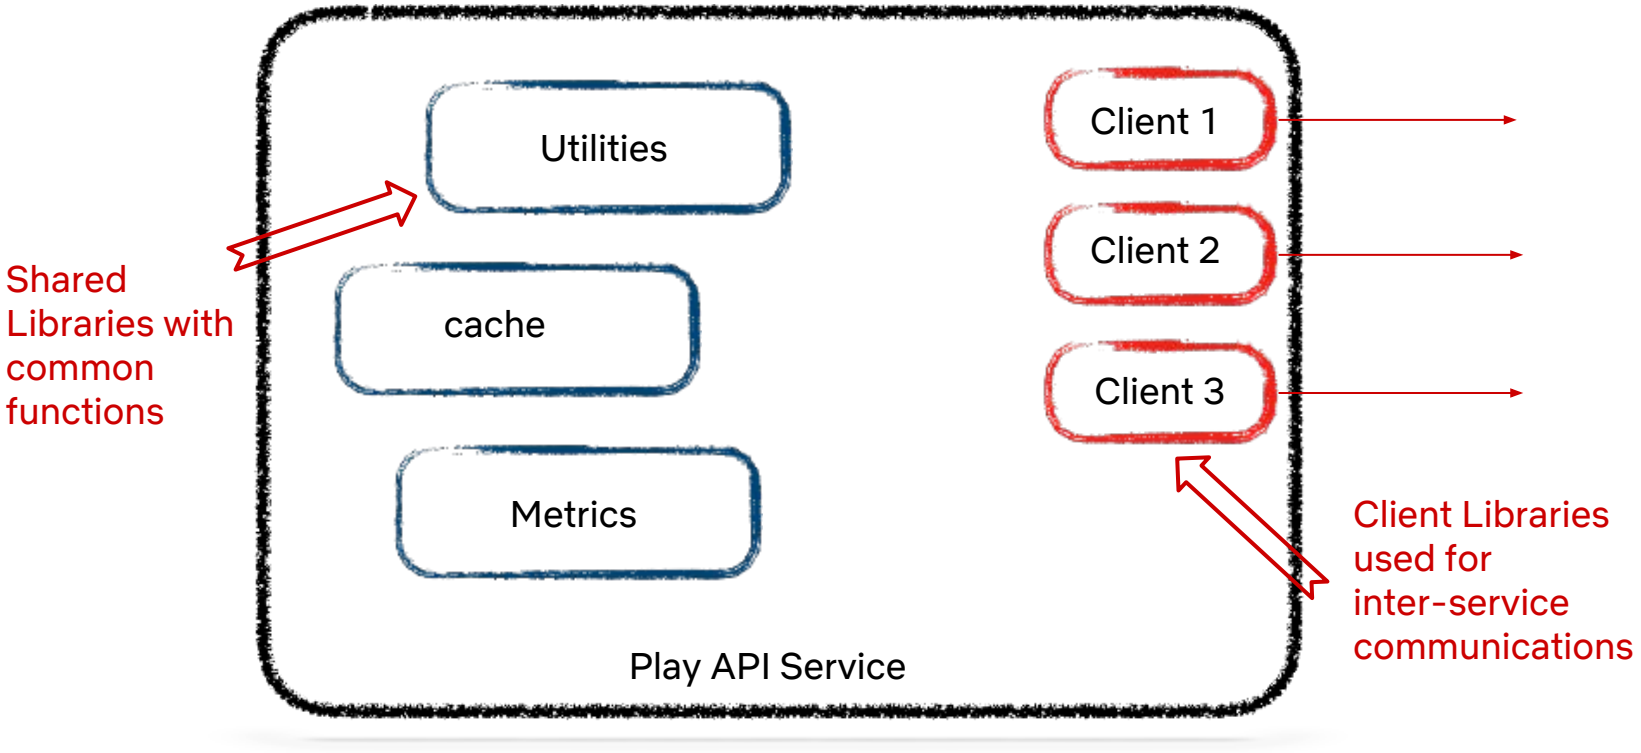

1) Binary Coupling

“Thick” shared libraries with 100s of dependent libraries (e.g. utilities jar)

Previous Architecture

Binary coupling => Distributed Monolith

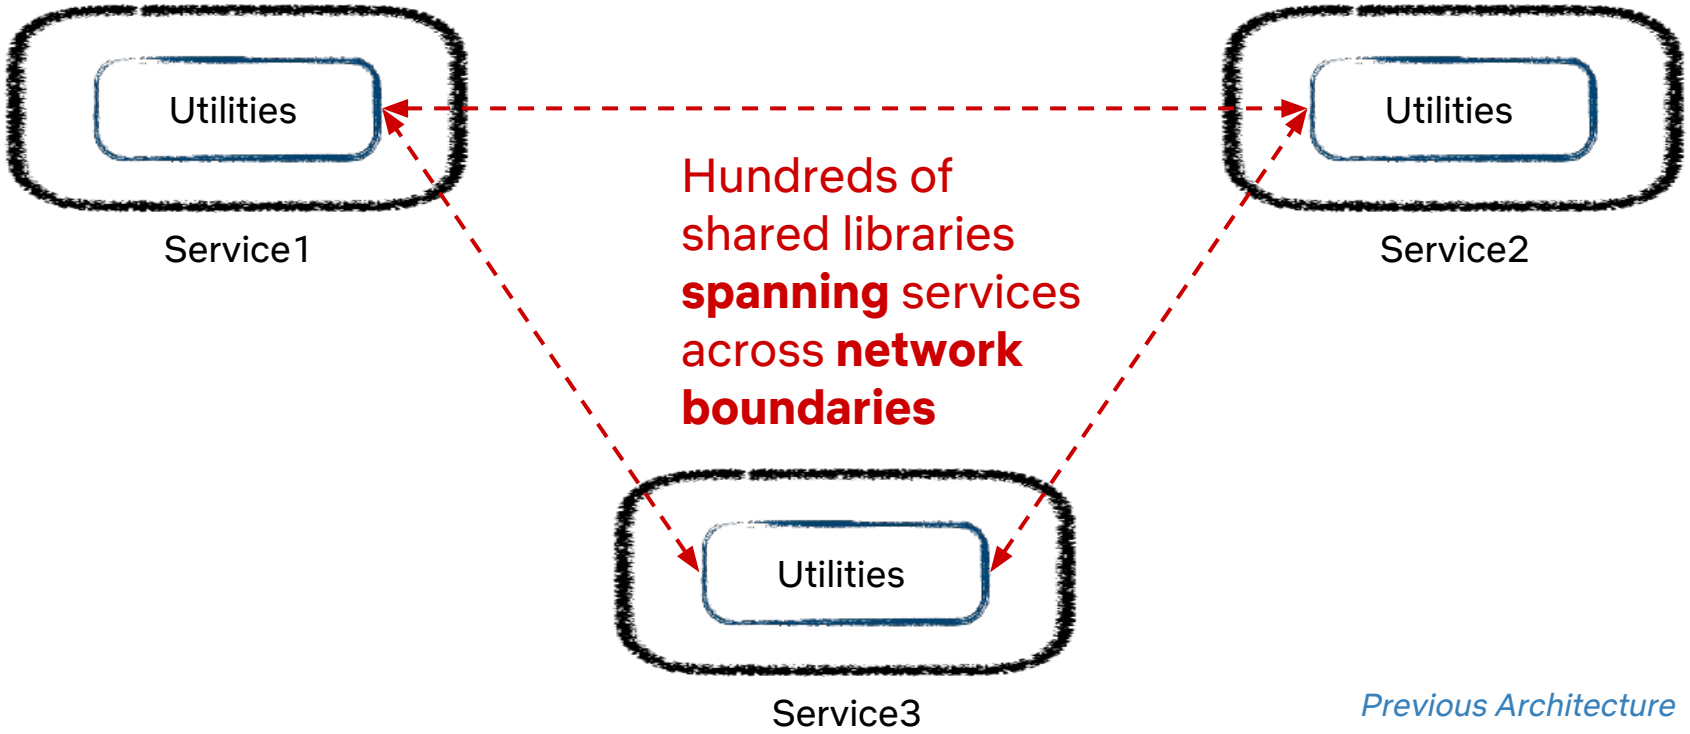

Previous Architecture

“The evils of too much coupling between services are far worse than the problems caused by code duplication”

- Sam Newman (Building Microservices)

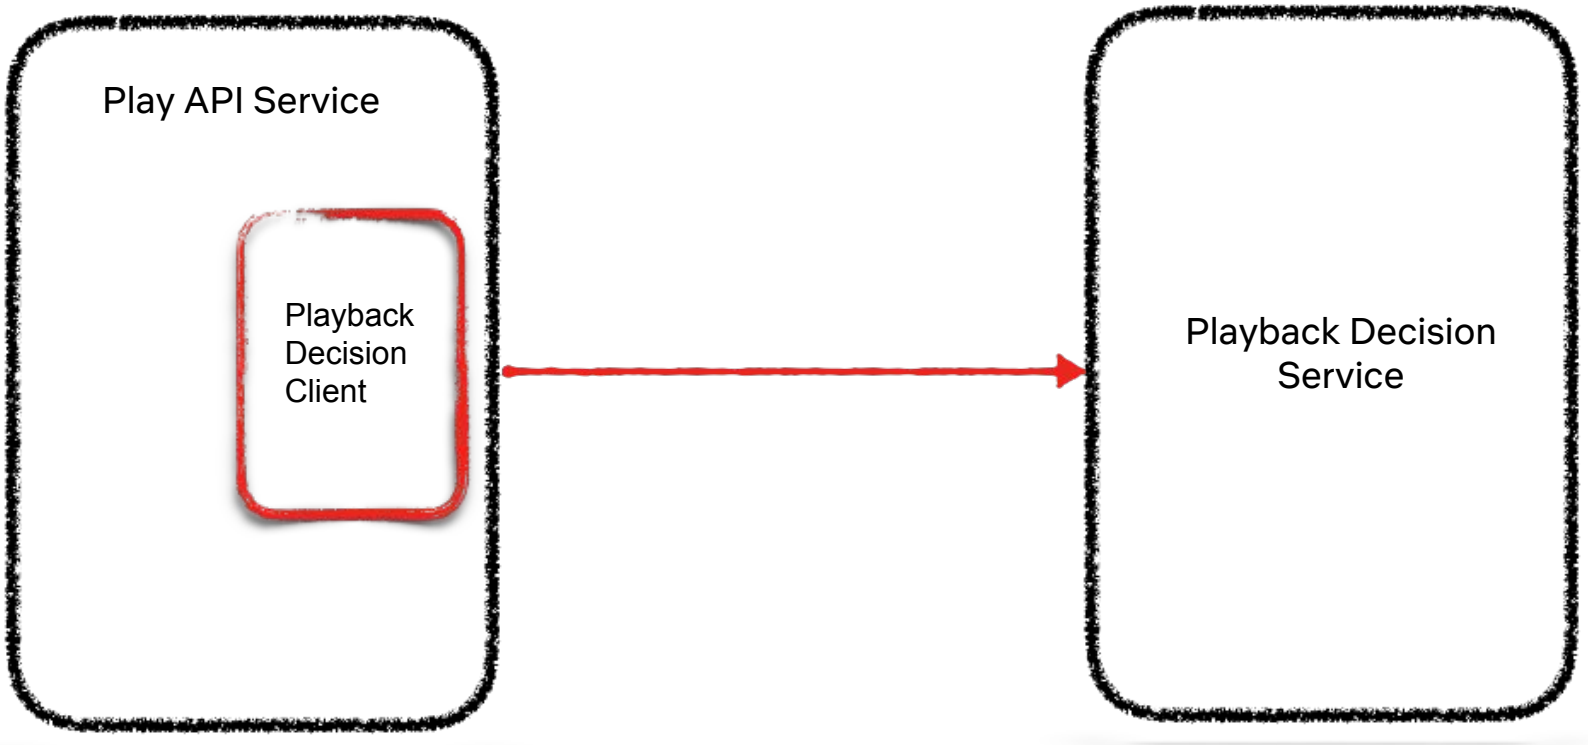

Previous Architecture

Clients with heavy Fallbacks

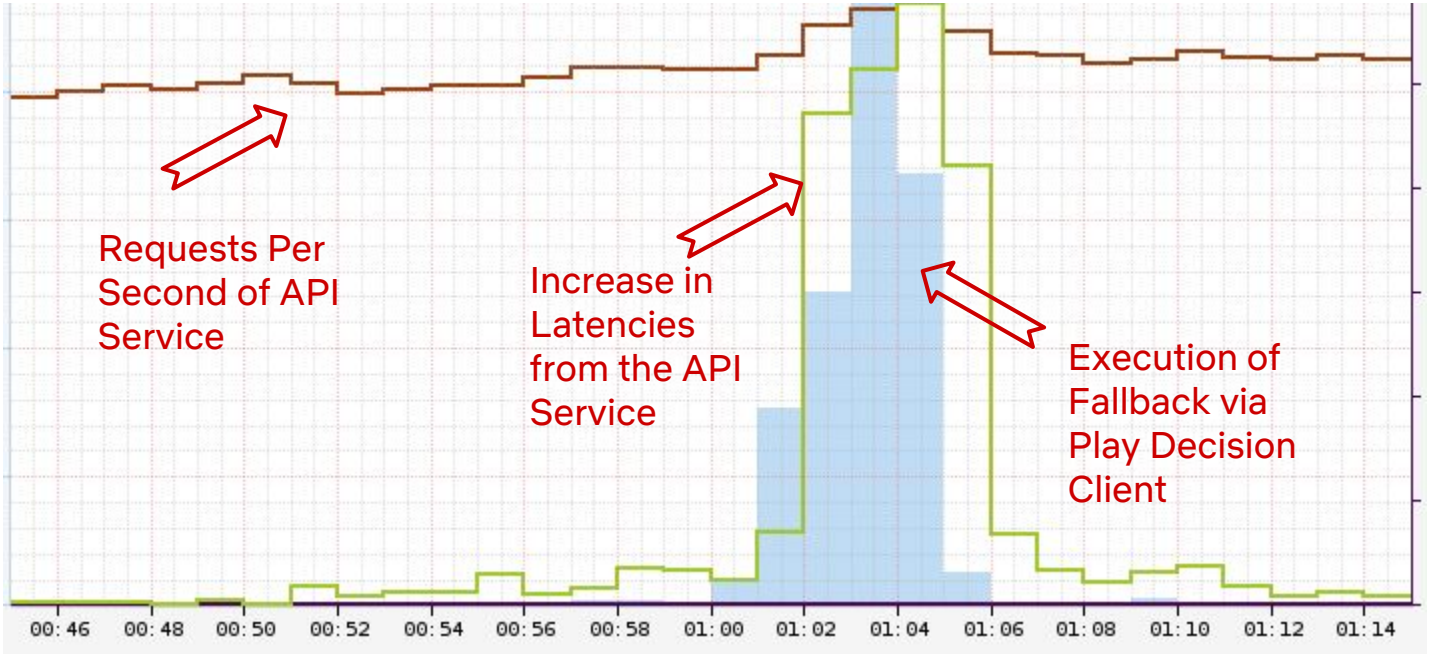

2) Operational Coupling

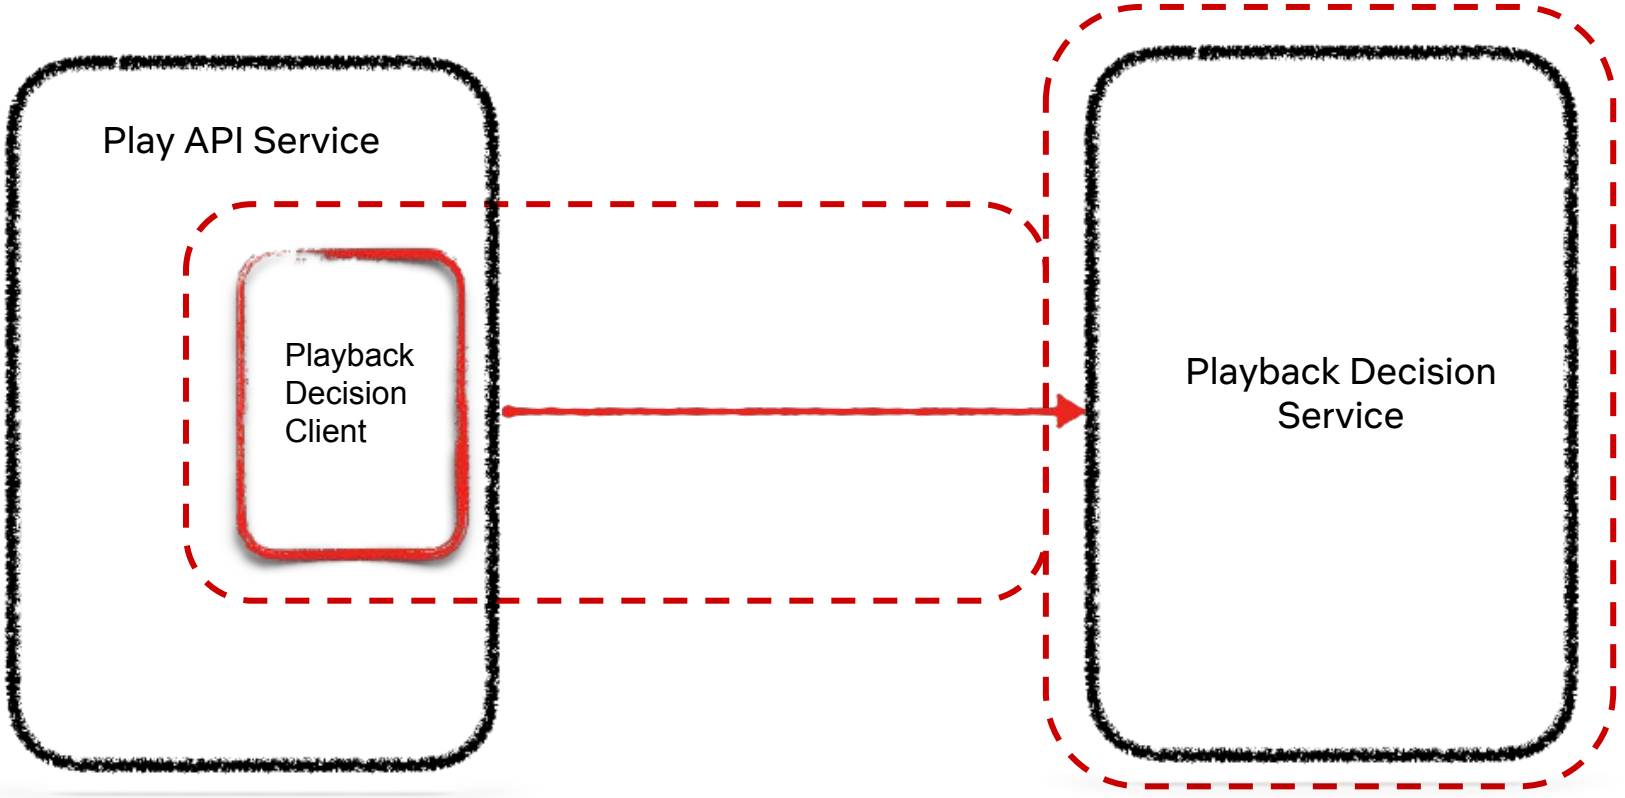

Previous Architecture

“Operational Coupling” might be an ok choice, if some services/teams are not yet ready to own and operate a highly available service.

Operational Coupling impacts Availability

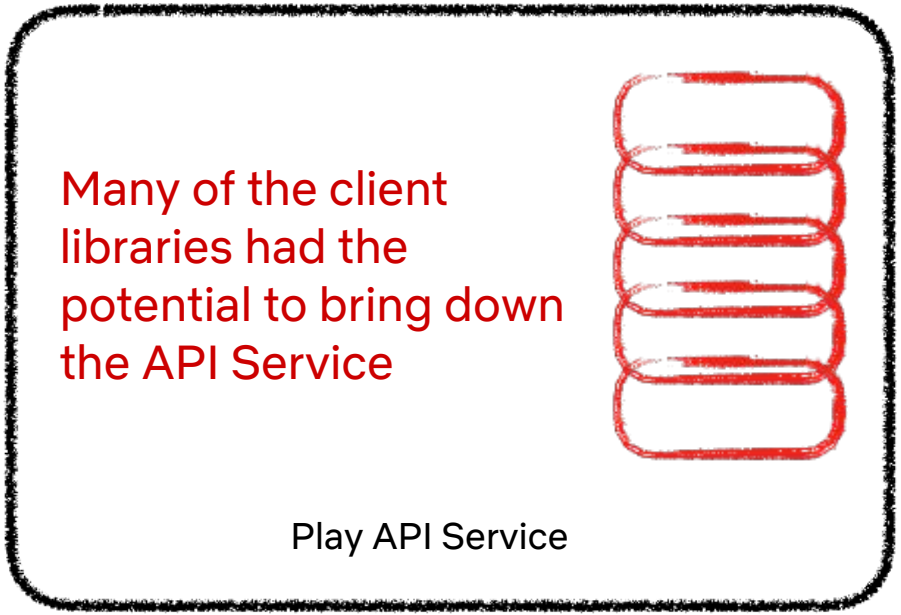

Previous Architecture

3) Language Coupling

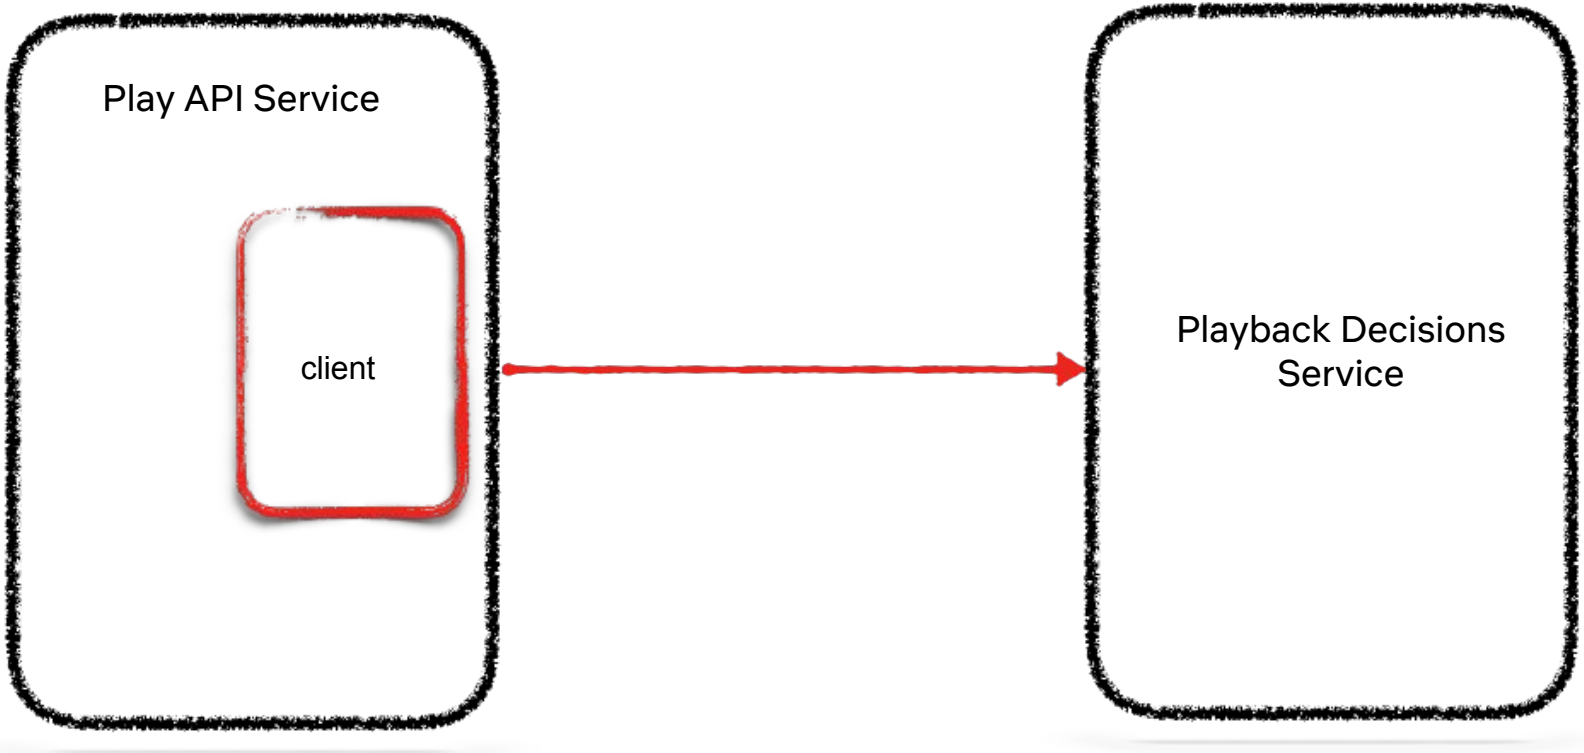

Previous Architecture

Communication Protocol

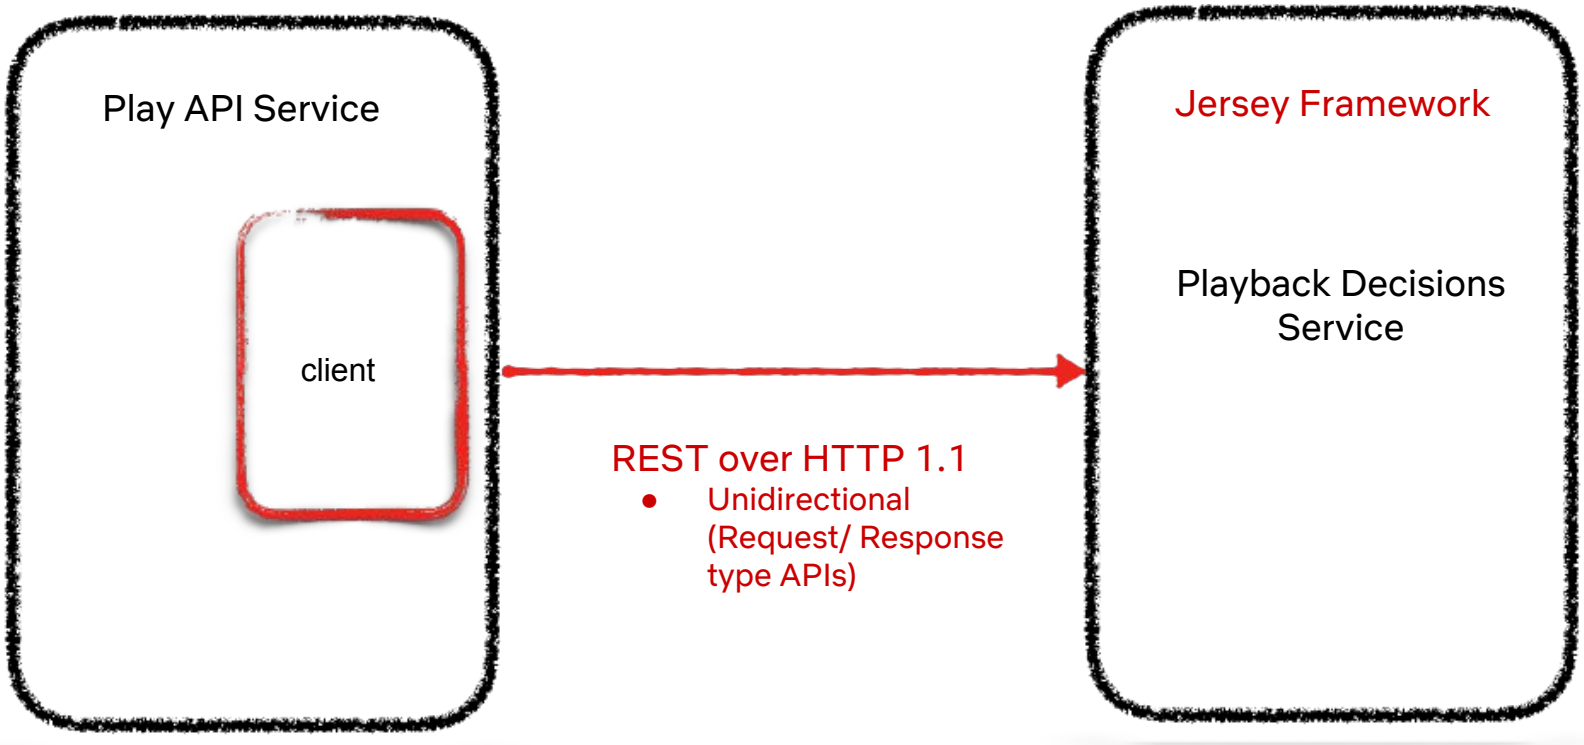

Previous Architecture

Requirements

Operationally “thin” Clients	No or limited shared libraries
Auto-generated clients for Polyglot support	Bi-Directional Communication

REST vs RPC

- At Netflix, most use-cases were modelled as Request/Response
  - REST was a simple and easy way of communicating between services; so choice of REST was more incidental rather than intentional
- Most of the services were not following RESTful principles.
  - The URL didn't represent a unique resource, instead the parameters passed in the call determined the response - effectively made them a RPC call
- So we were agnostic to REST vs RPC as long as it meets our requirements


GRPC

Previous Architecture

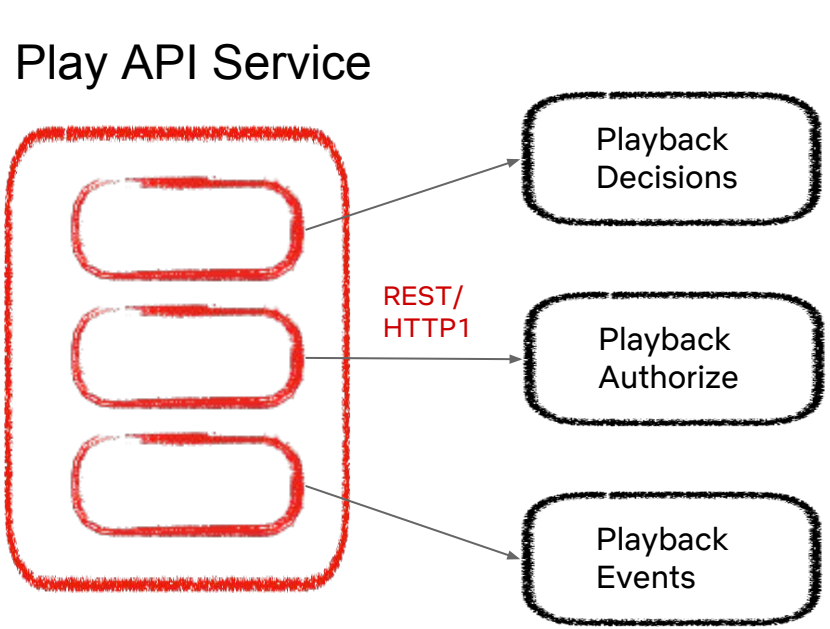

1) Operationally Coupled Clients
2) High Binary Coupling
3) Only Java
4) Unidirectional communication

Current Architecture

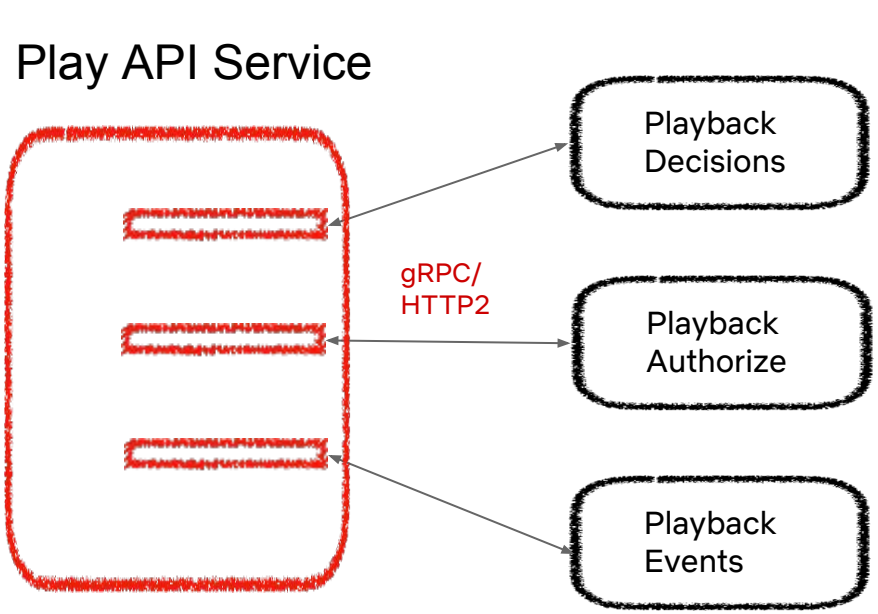

1) Minimal Operational Coupling
2) Limited Binary Coupling
3) Beyond Java
4) Beyond Request/ Response

Type 1 Decision: Appropriate Coupling

Consider “thin” auto-generated clients with bi-directional communication and minimize code reuse across service boundaries

Three Type 1 Decisions to Consider

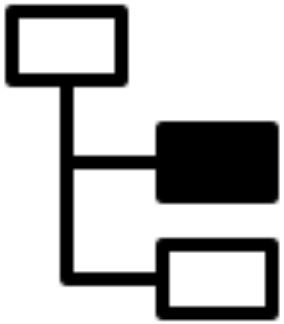

Appropriate Coupling

Synchronous vs Asynchronous

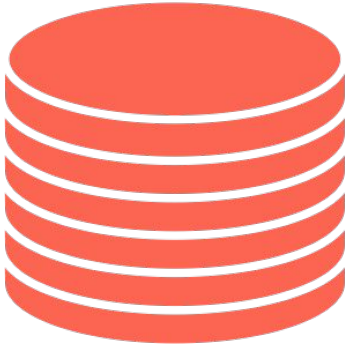

Data Architecture

```
PlayData getPlayData(string customerId, string titleId,
string deviceId){
    CustomerInfo custInfo = getCustomerInfo(customerId);
    DeviceInfo deviceInfo = getDeviceInfo(deviceId);
    PlayData playdata = decidePlayData(custInfo,
deviceInfo, titleId);
    return playdata;
}
```

Typical Synchronous Architecture

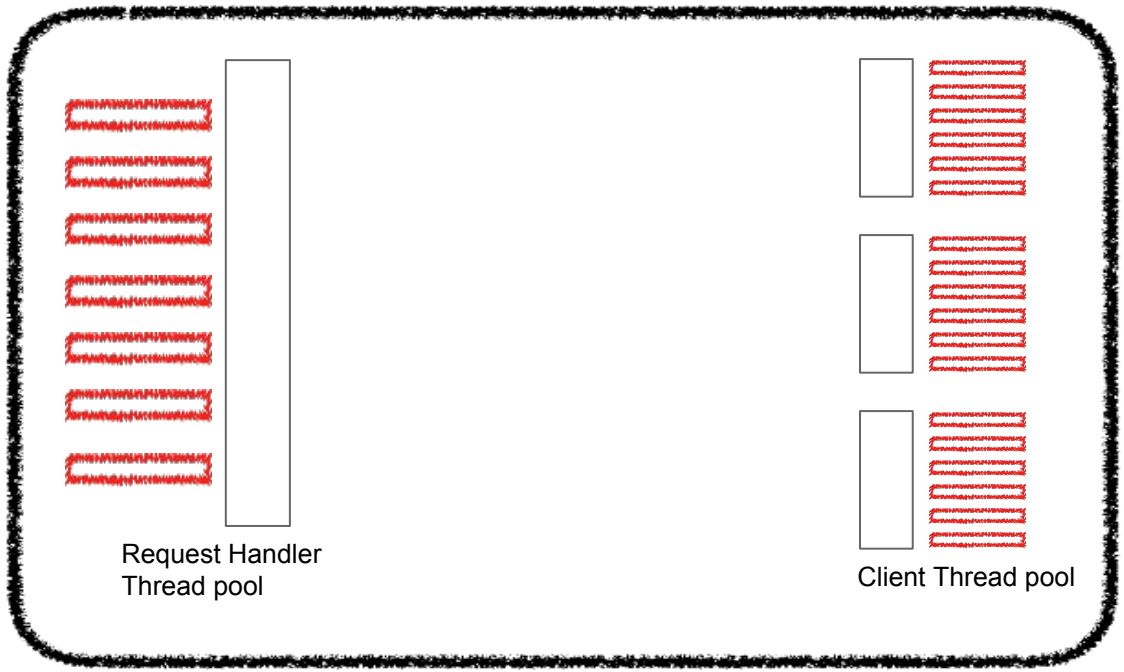

Typical Synchronous Architecture

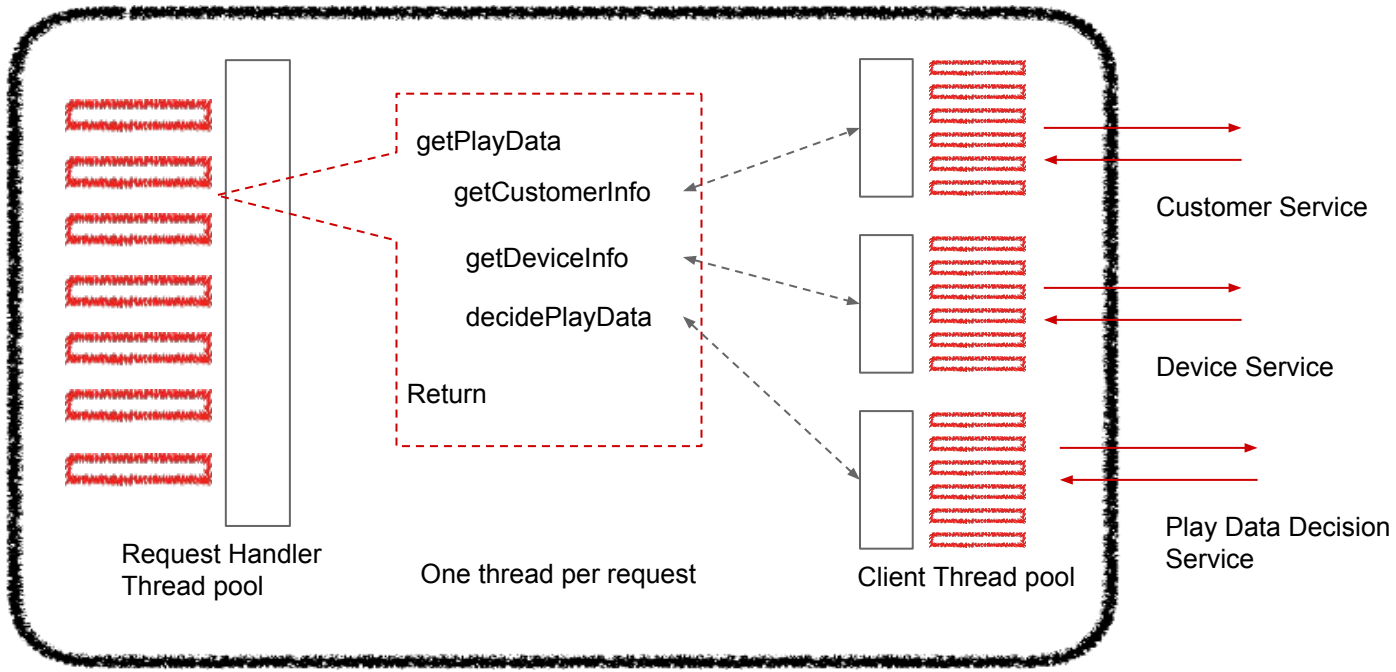

Typical Synchronous Architecture

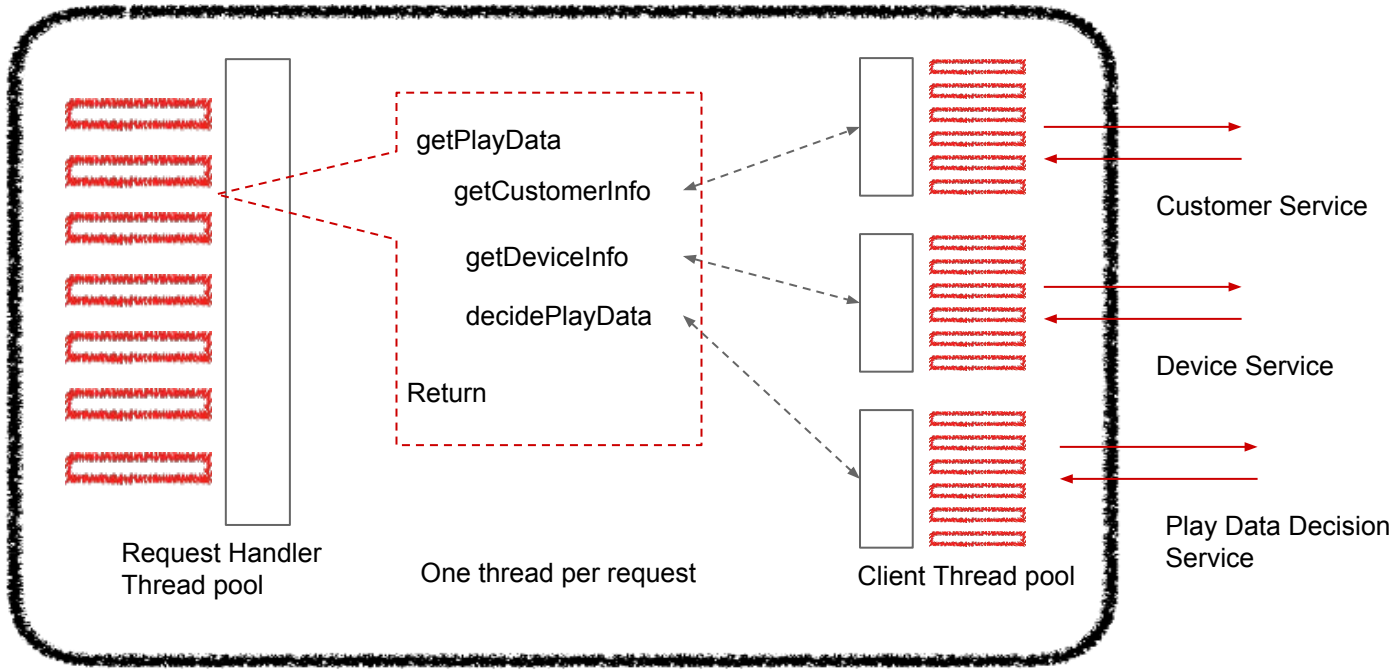

Blocking Request Handler

Blocking Client I/O

Typical Synchronous Architecture

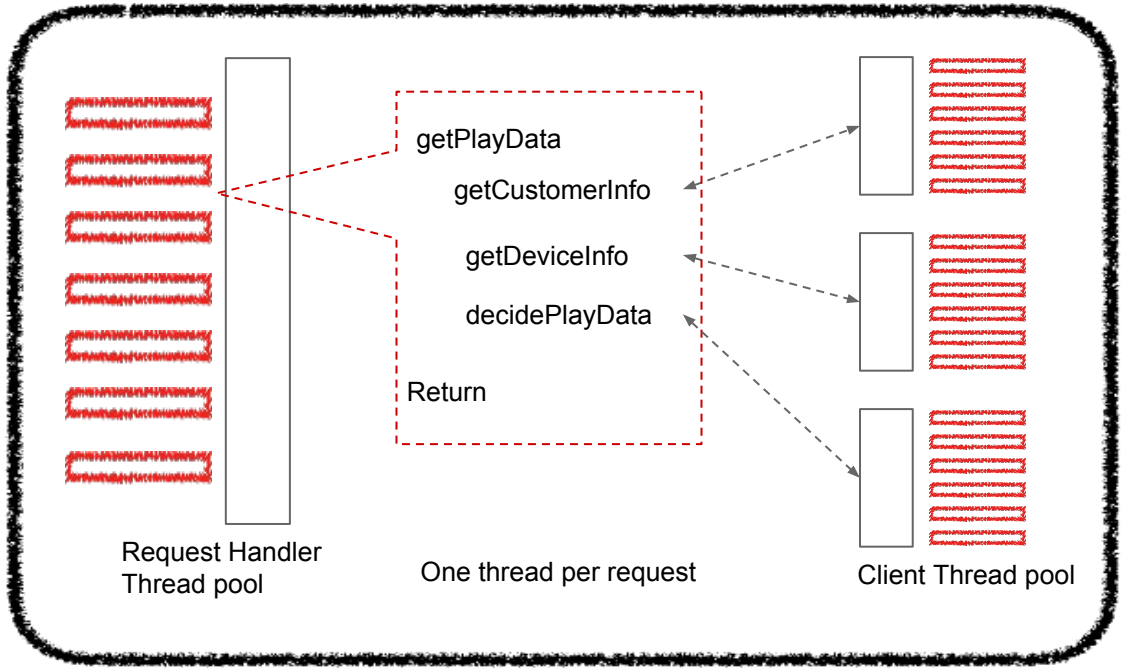

- Works for Simple Request/Response
- Works for Limited Clients

Blocking Request Handler

Blocking Client I/O

Beyond Request/Response

One Request - One Response

Request Play-data for Title X
Receive Play-data for Title X

One Request - Stream Response

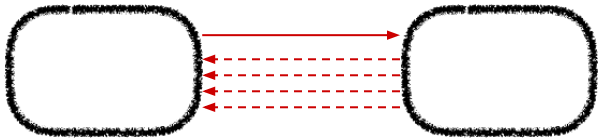

Request Play-data for Titles X,Y,Z
Receive Play-data for Title X
Receive Play-data for Title Y
Receive Play-data for Title Z

Stream Request - One Response

Request Play-data for Title X
Request Play-data for Title Y
Request Play-data for Title Z
Receive Play-data for Titles X,Y,Z

Stream Request - Stream Response

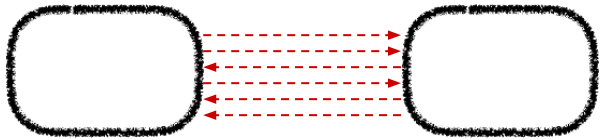

Request Play-data for Title X
Request Play-data for Title Y
Receive Play-data for Title X
Get Play-data for Title Z
Receive Play-data for Title Y
Receive Play-data for Title Z

Asynchronous Architecture

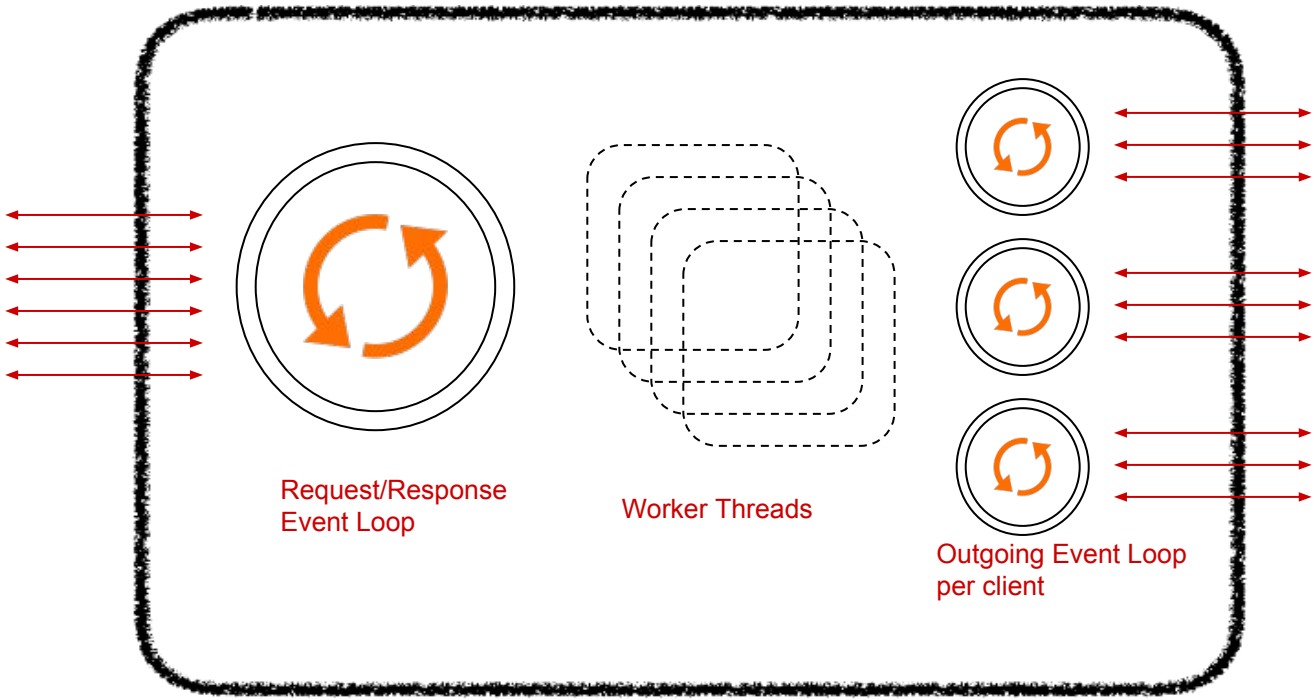

```
PlayData getPlayData(string customerId, string titleId,
string deviceId){
    Zip(getCustomerInfo(customerId),
         getDeviceInfo(deviceId),
         (custInfo, deviceInfo) ->
           return decidePlayData(custInfo, deviceInfo,
        titleId)
        );
}
```

Asynchronous Architecture

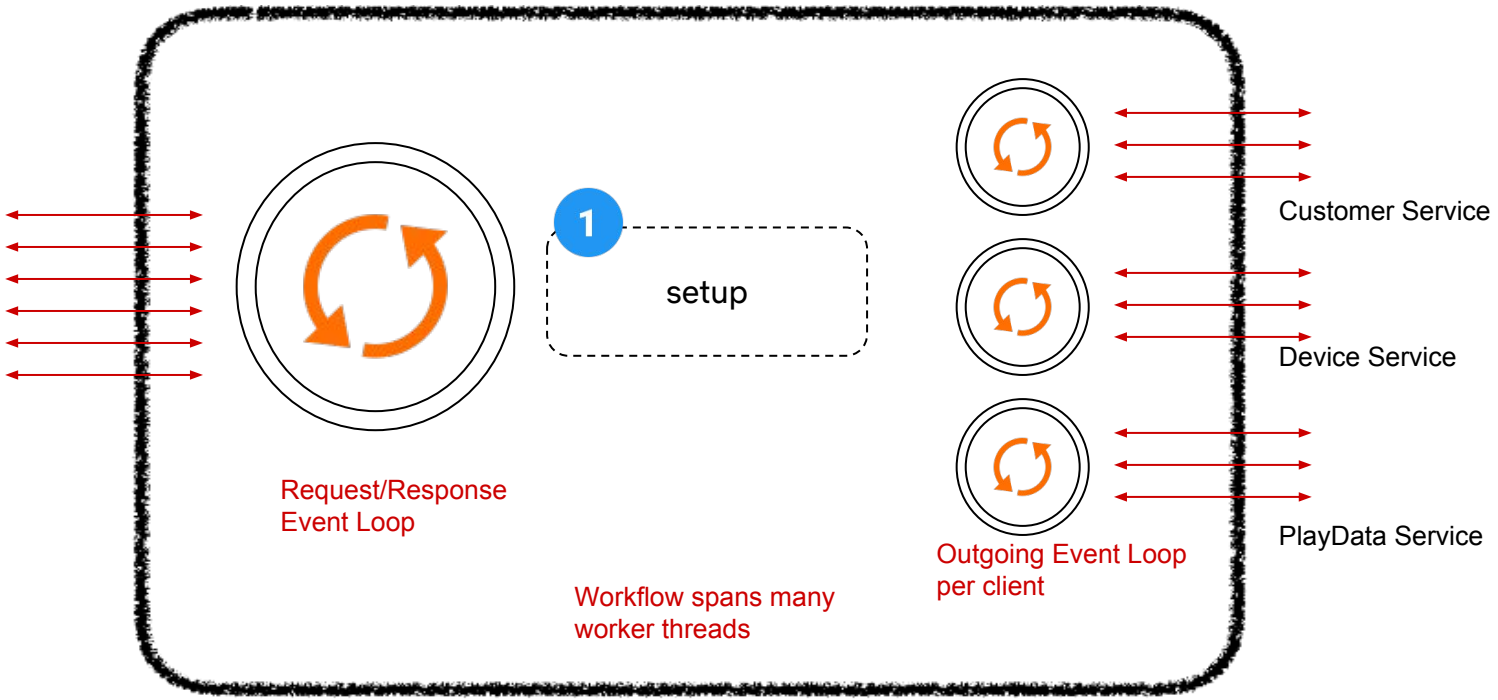

Asynchronous Architecture

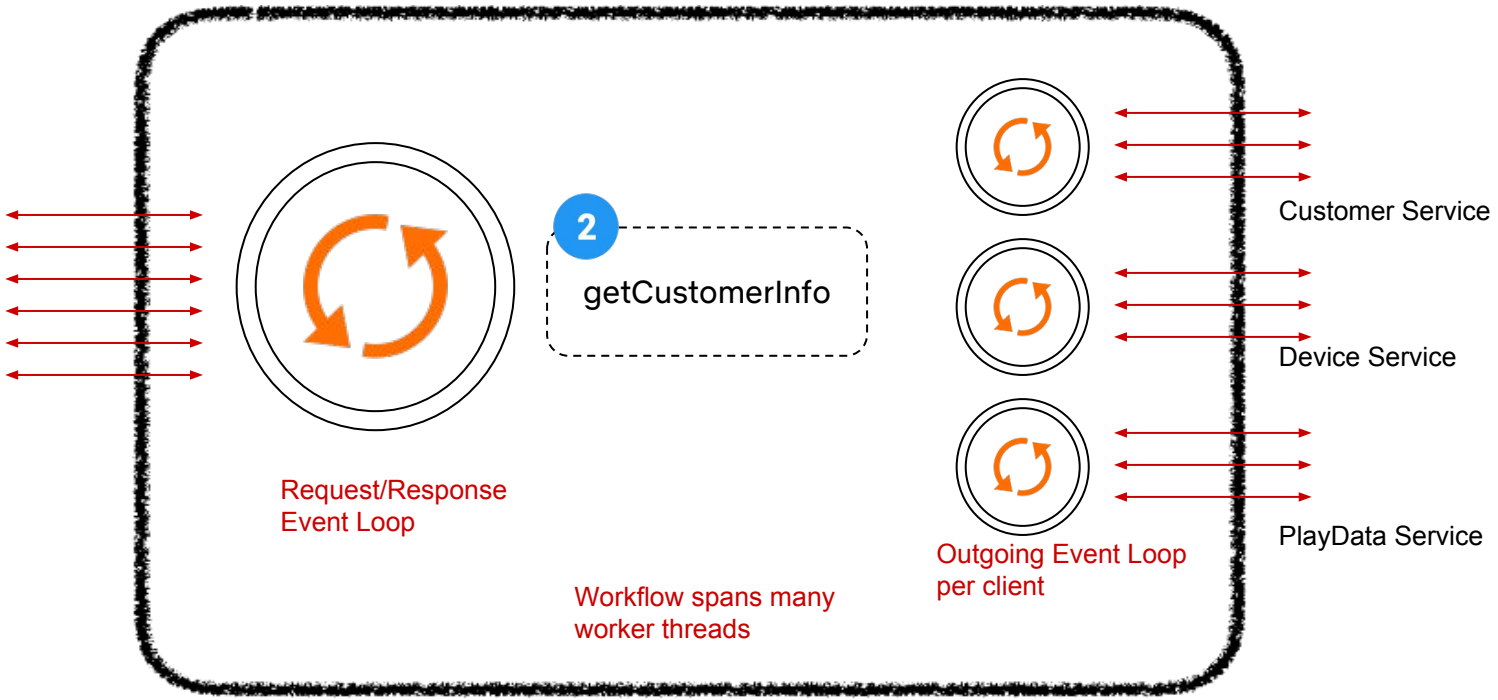

Asynchronous Architecture

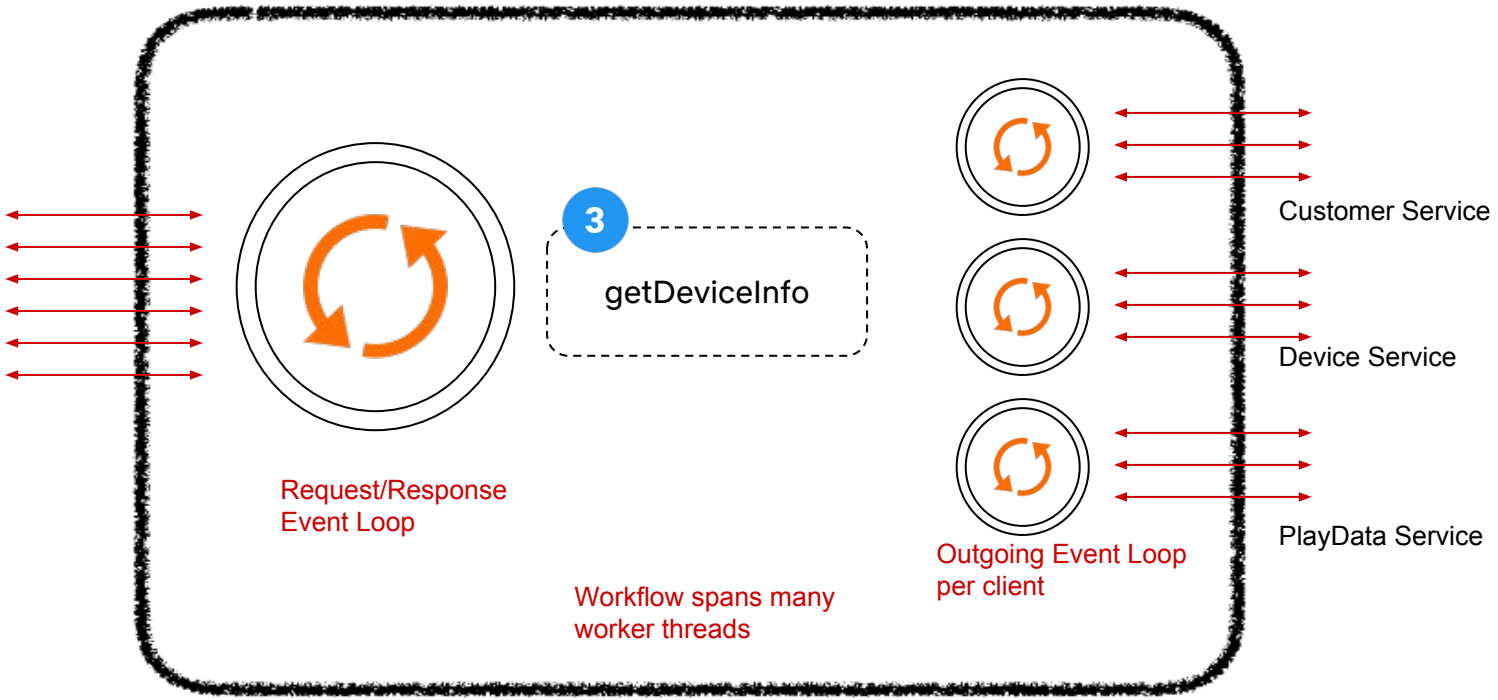

Asynchronous Architecture

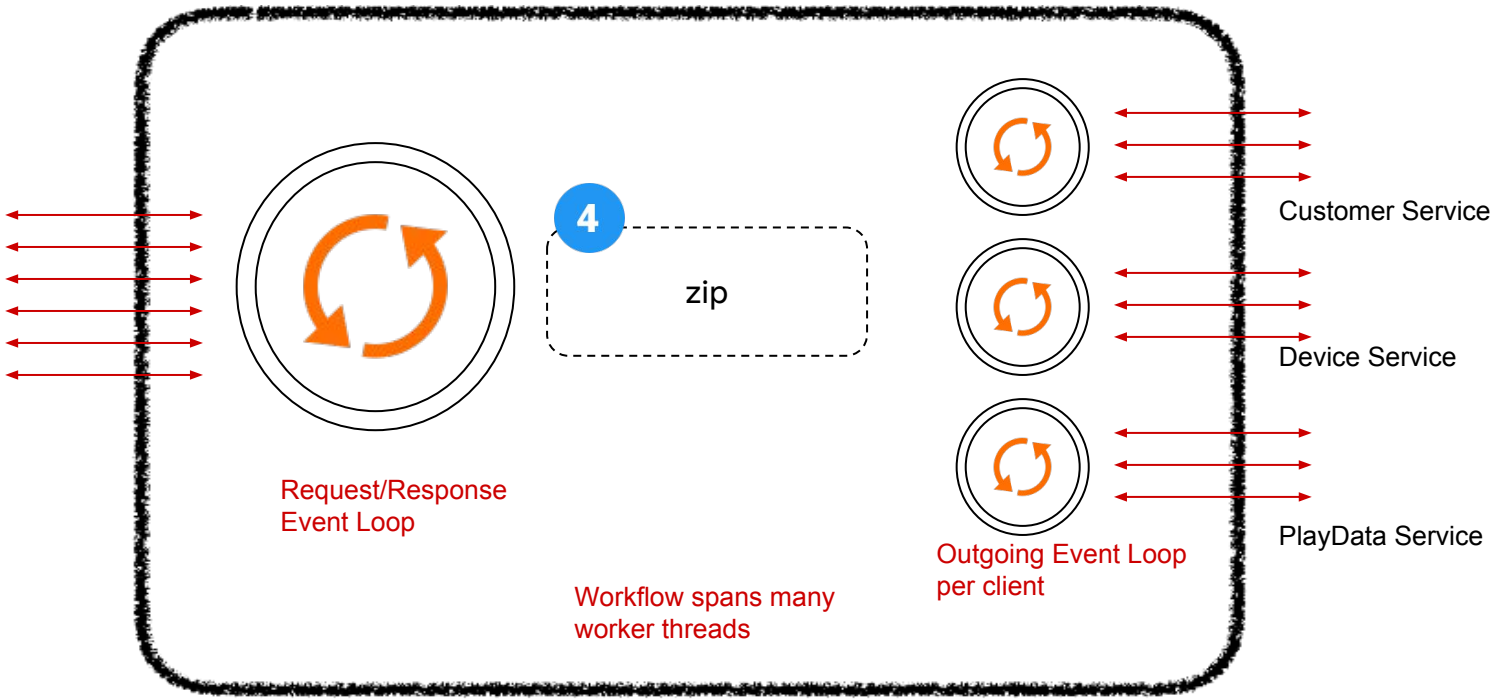

Asynchronous Architecture

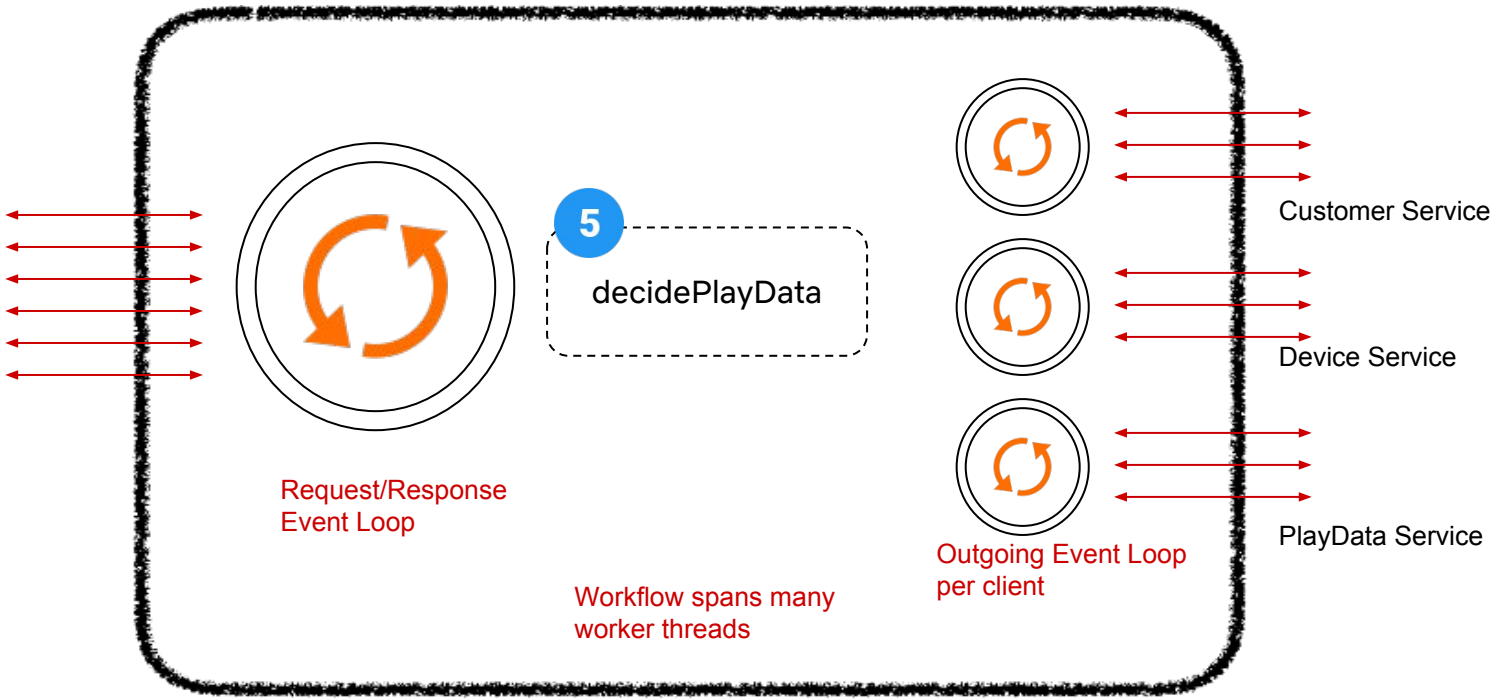

Workflow spans multiple threads

- All context is passed as messages from one processing unit to another.
- If we need to follow and reason about a request, we need to build tools to capture and reassemble the order of execution units
- None of the calls can block

Asynchronous Architecture

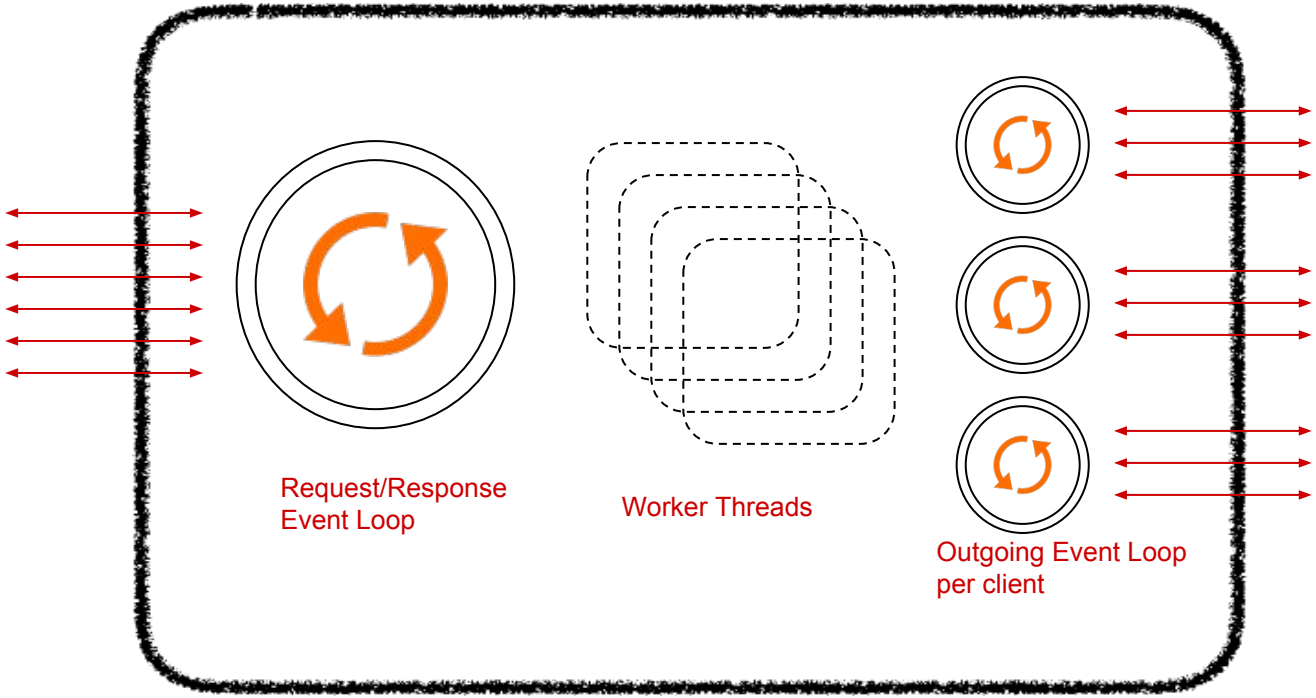

Asynchronous Request Handler **Non-Blocking I/O**

Synchrony

Ask: Do you really have a need beyond Request/Response?

Synchronous Execution + Asynchronous I/O

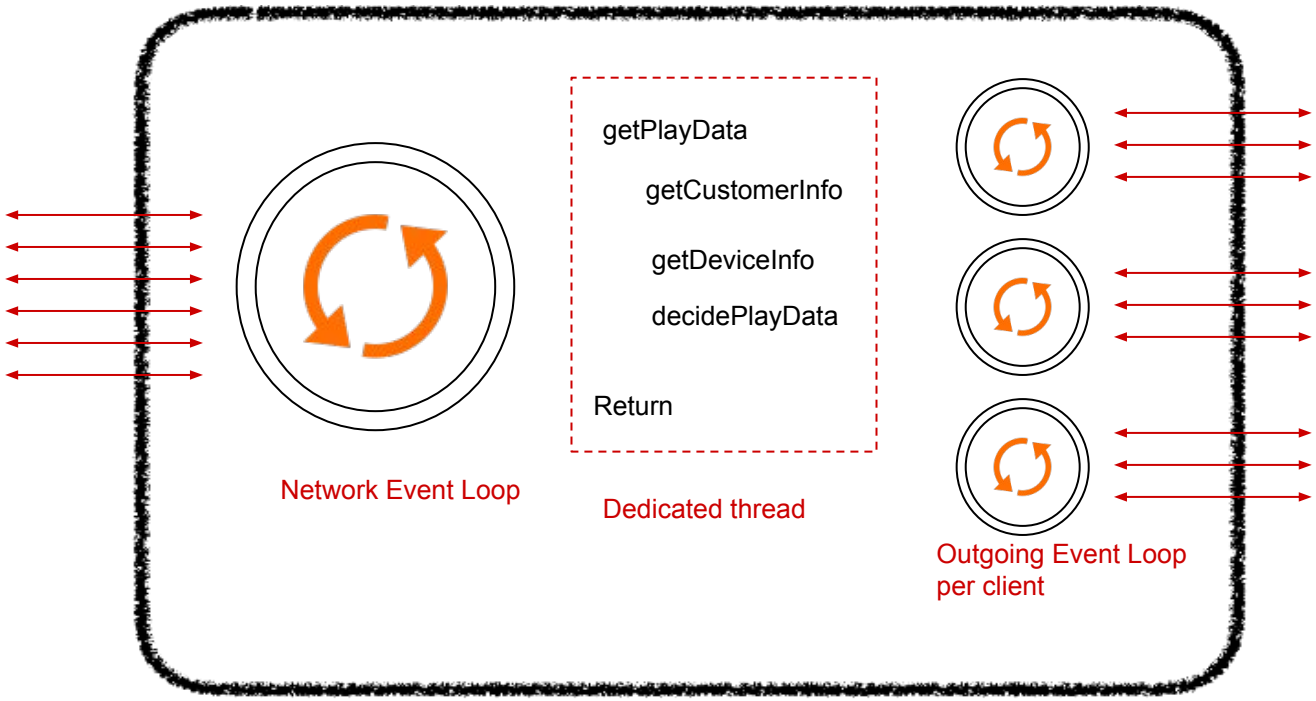

Current Architecture

Type 1 Decision: Synchronous vs Asynchronous

If most of your APIs fit the Request/Response pattern, consider a synchronous request handler, with nonblocking I/O

Three Type 1 Decisions to Consider

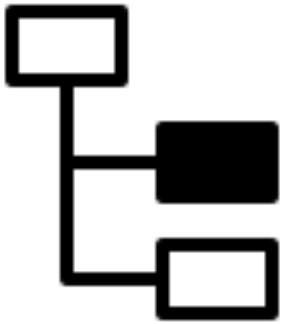

- Appropriate Coupling

- Synchronous vs Asynchronous

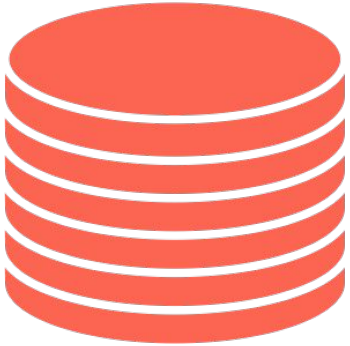

Data Architecture

Without an intentional Data Architecture, Data becomes its own monolith

What a Data Monolith looks like

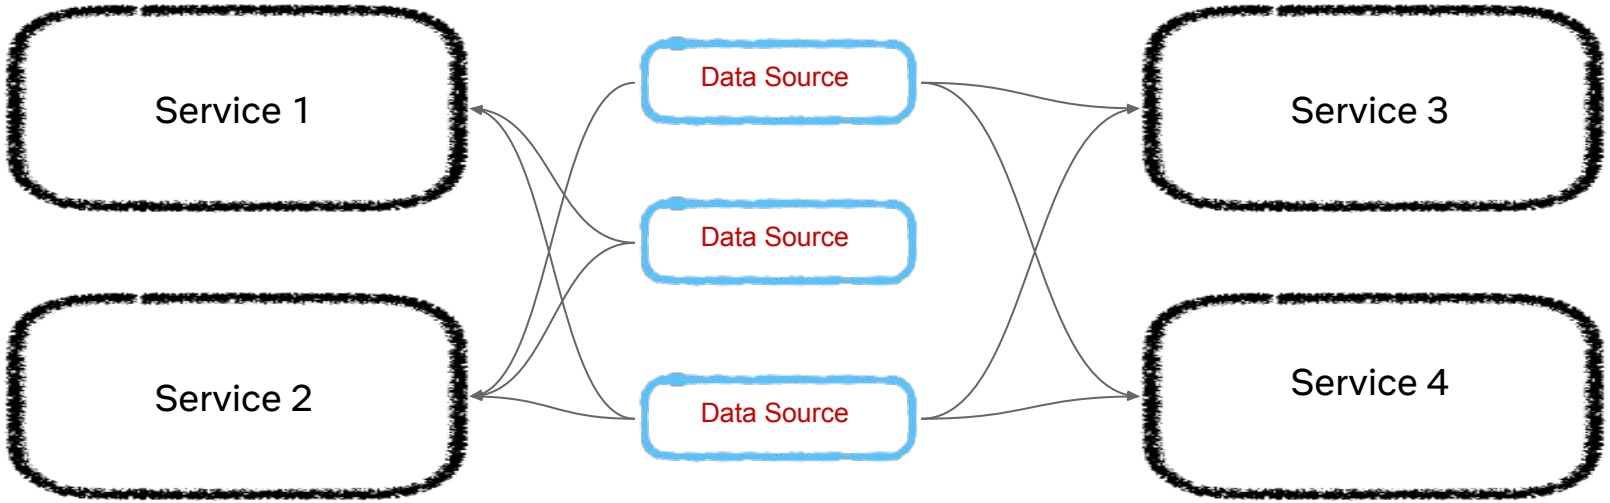

Previous Architecture

What a Data Monolith looks like

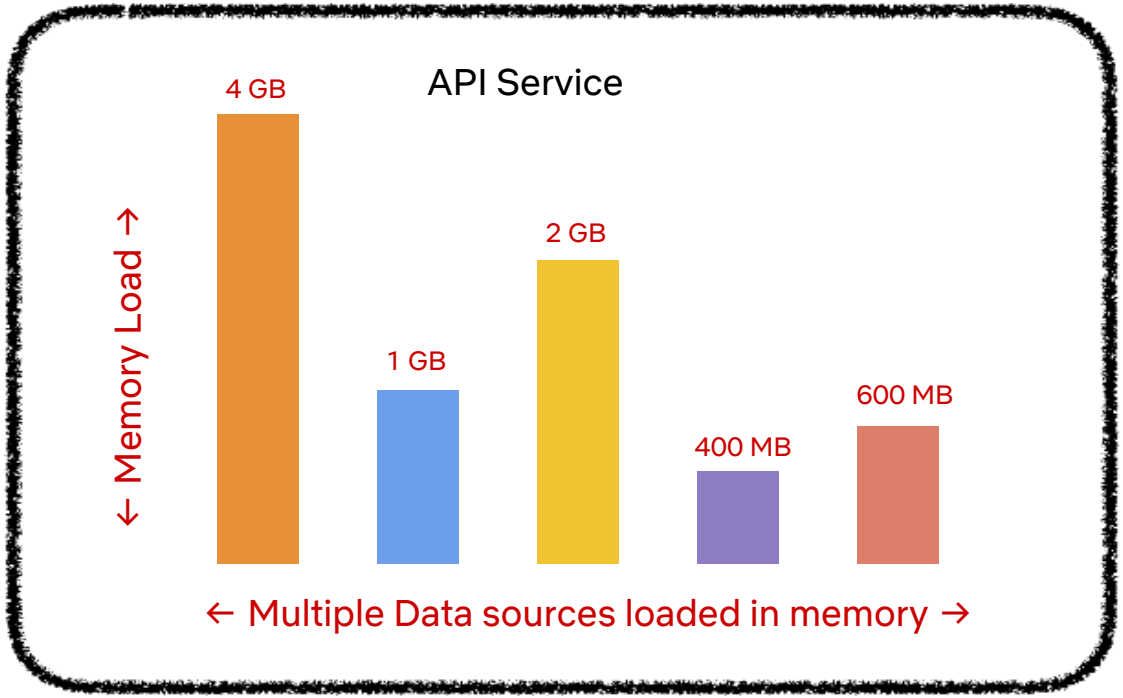

Previous Architecture

What a Data Monolith looks like

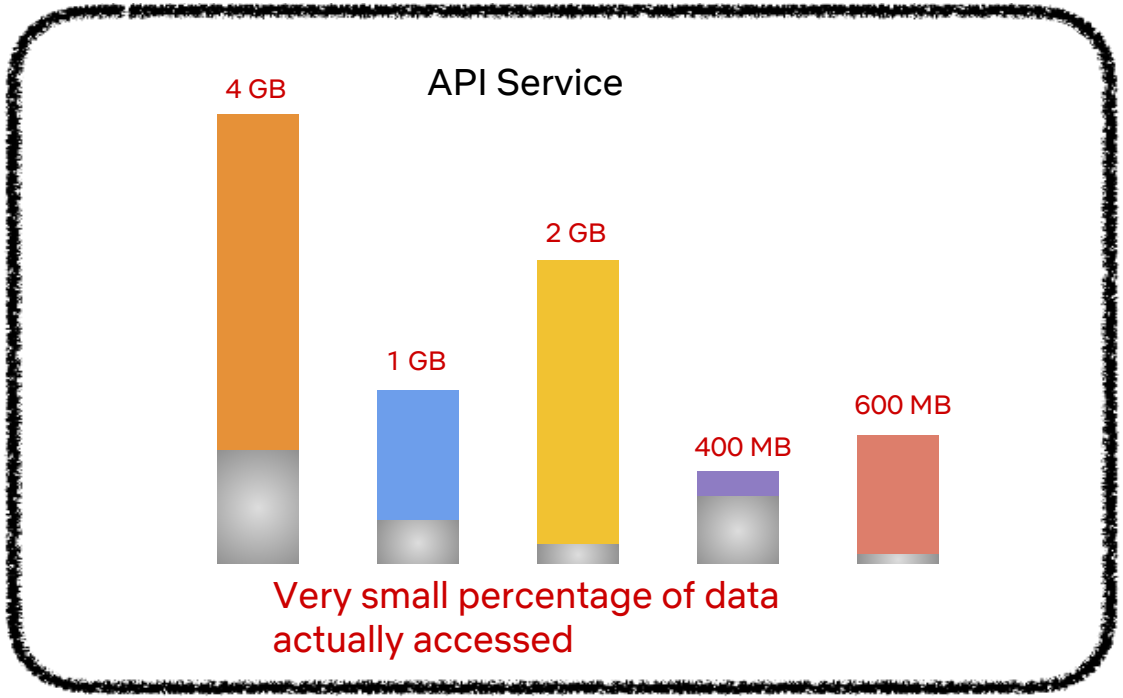

Previous Architecture

What a Data Monolith looks like

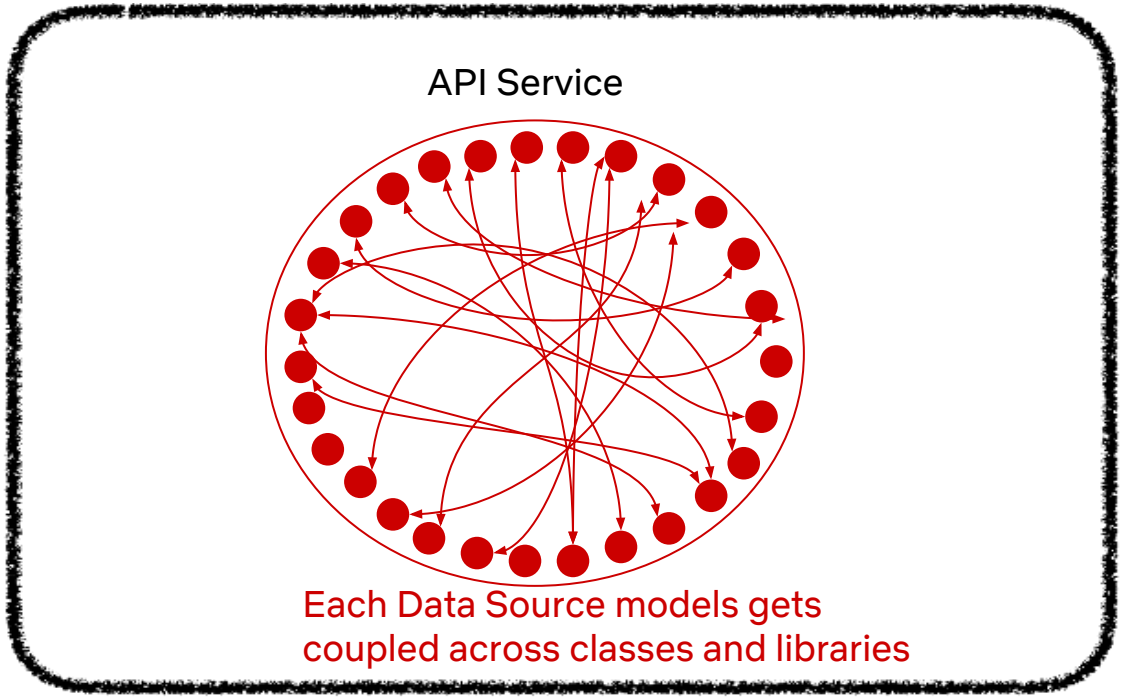

Previous Architecture

What a Data Monolith looks like

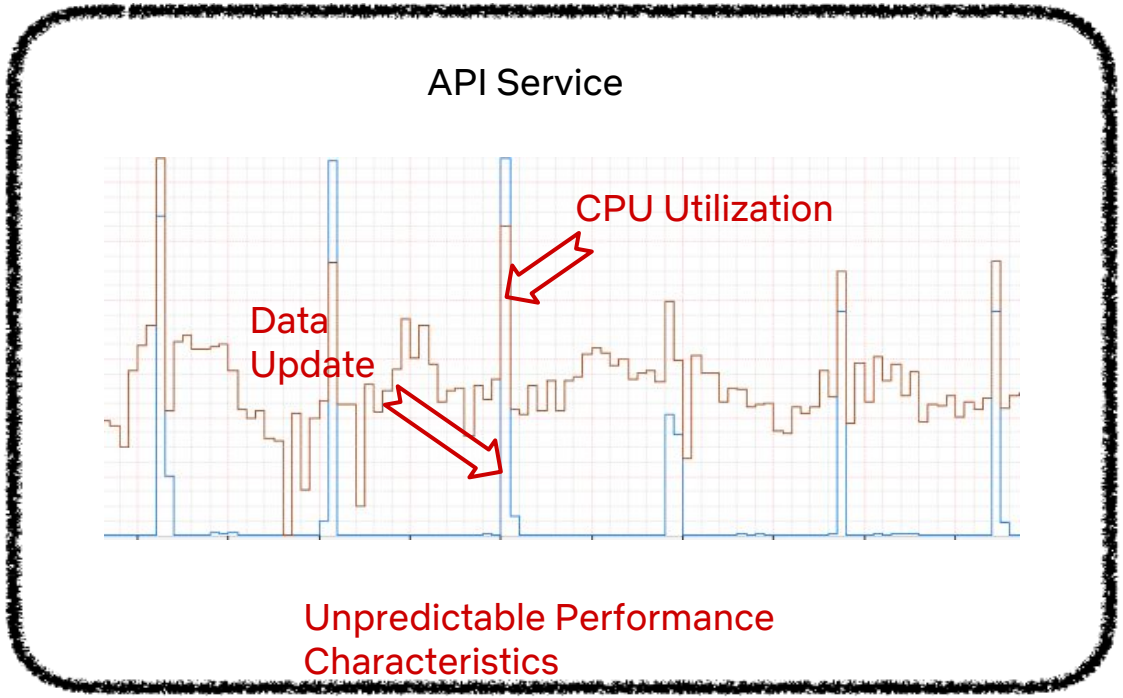

Previous Architecture

What a Data Monolith looks like

API Service

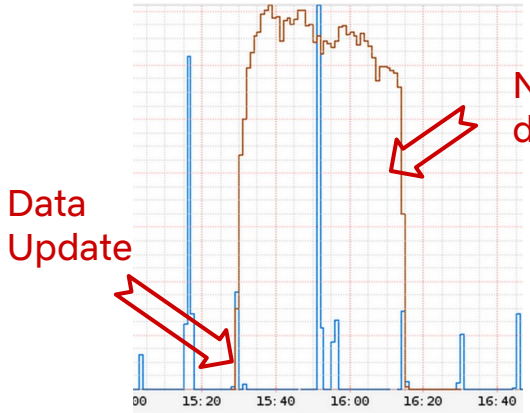

Potential to bring down the service

Previous Architecture

"All problems in computer science can be solved by another level of indirection."

David Wheeler

(World's first Comp Sci PhD)

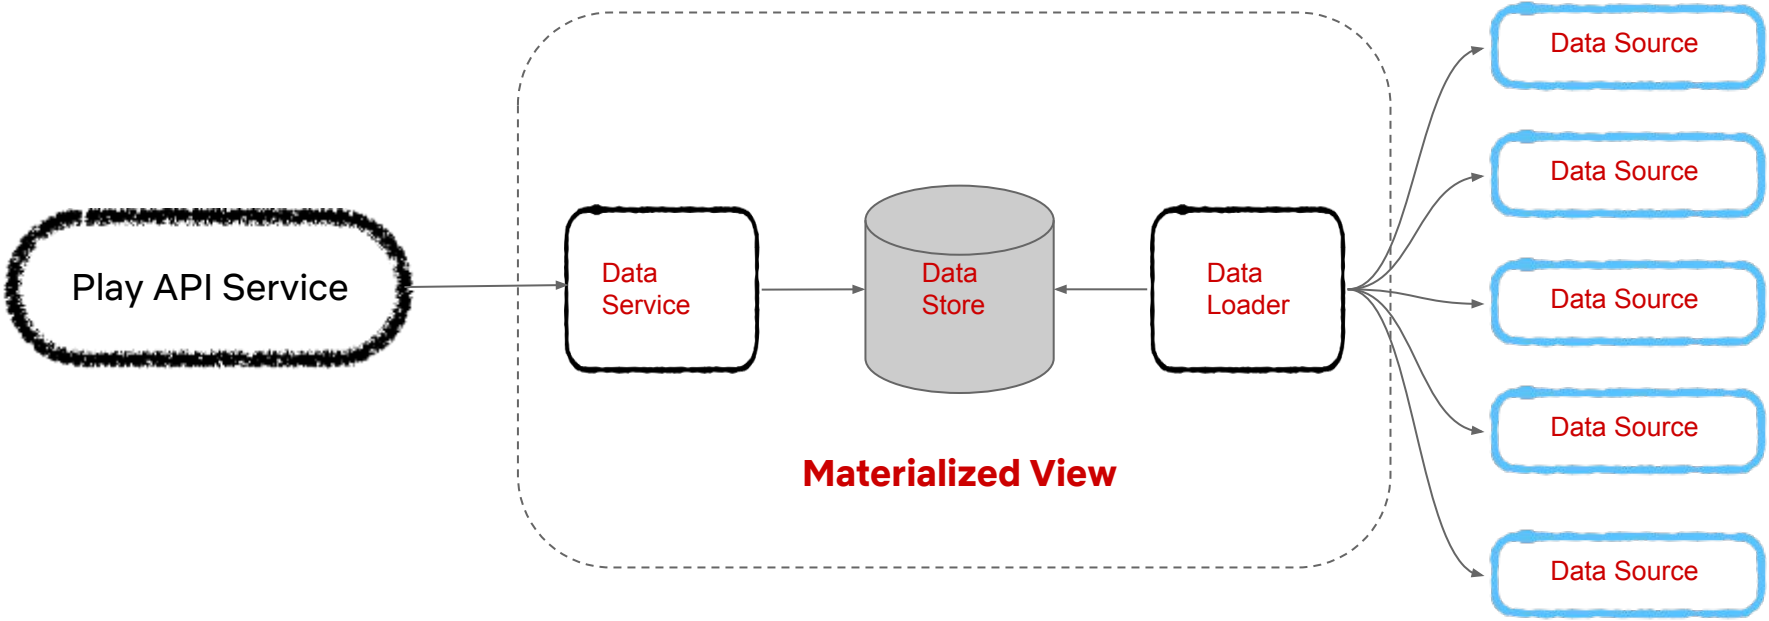

Current Architecture

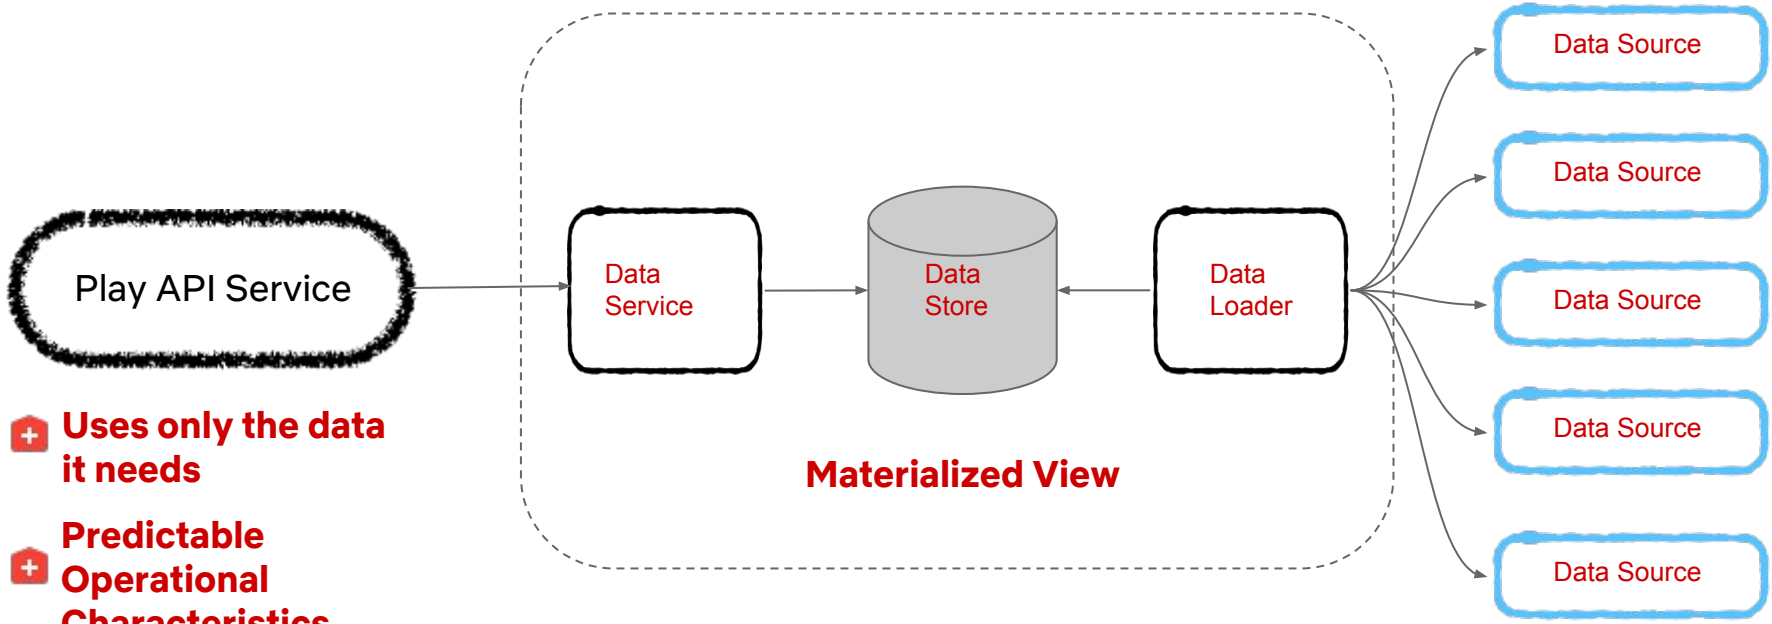

- **Uses only the data it needs**
- **Predictable Operational Characteristics**
- **Reduced Dependency chain**

Current Architecture

Type 1 Decision: Data Architecture

Isolate Data from the Service. At the very least, ensure that data sources are accessed via a layer of abstraction, so that it leaves room for extension later

Three Type 1 Decisions to Consider

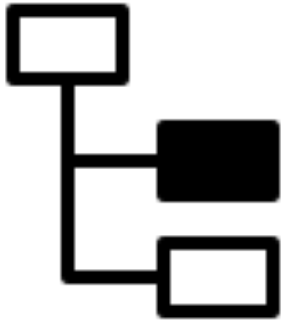

- Appropriate Coupling

- Synchrony

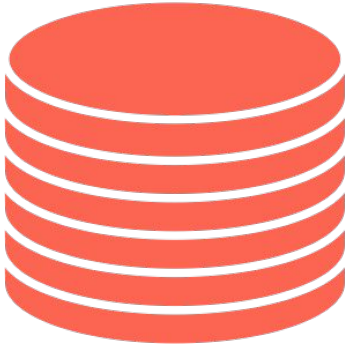

- Data Architecture

For Type 2 decisions, choose a path,
experiment and iterate

Guiding Principle: Identify your Type 1 and Type 2 decisions; Spend 80% of your time debating and aligning on Type 1 Decisions

Identity

Type 1/2 Decisions

Evolvability

NETFLIX

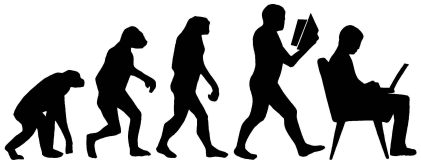

An Evolutionary Architecture supports guided and incremental change as first principle among multiple dimensions

- ThoughtWorks

Choosing a microservices architecture with appropriate coupling allows us to evolve across multiple dimensions

How evolvable are the Type 1 decisions

Known Unknowns

Change Play API	Previous Architecture	Current Architecture
Asynchronous?		
Polyglot services?		
Bidirectional APIs?		
Additional Data Sources?		

Potential Type 1 decisions in the future?

Change Play API	Previous Architecture	Current Architecture
Containers?		?
Serverless?		?

And we fully expect that there will be Unknown Unknowns

As we evolve, how to ensure we are not breaking our original goals?

Use Fitness Functions to guide change

High Availability

Low Latency

Evolvability

Simplicity

Scalable

Reliability

Continuous Integration

High Throughput

Observability

Developer Productivity

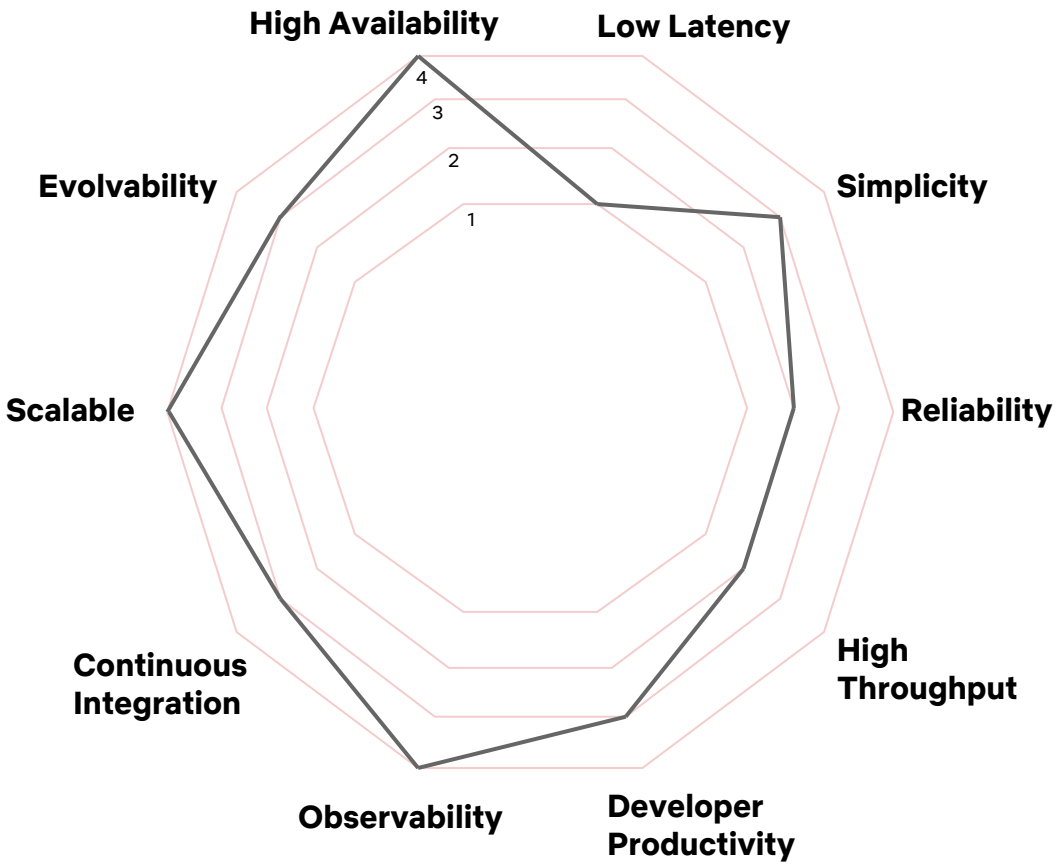
High Availability
Low Latency
Simplicity
Reliability
High Throughput
Developer Productivity
Observability
Continuous Integration
Scalable
Evolvability
4
3
2
1

Why Simplicity over Reliability?

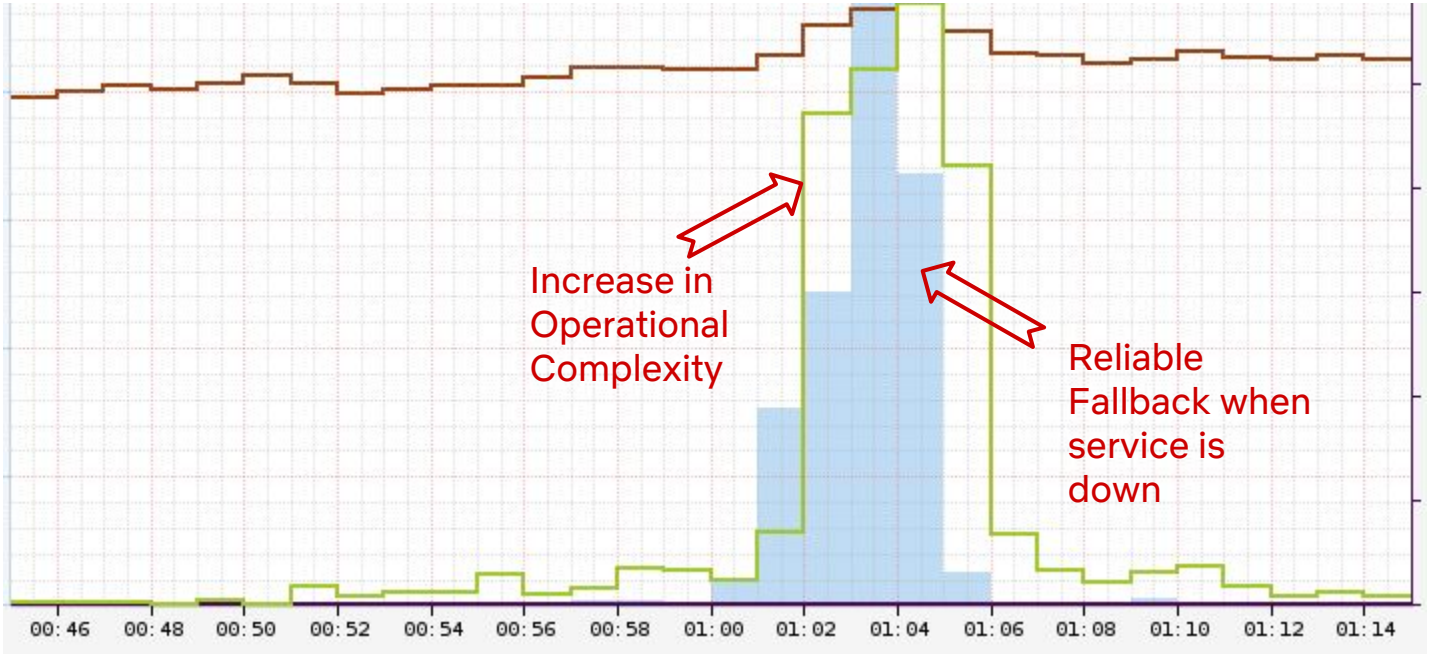

Why Scalability over Throughput?

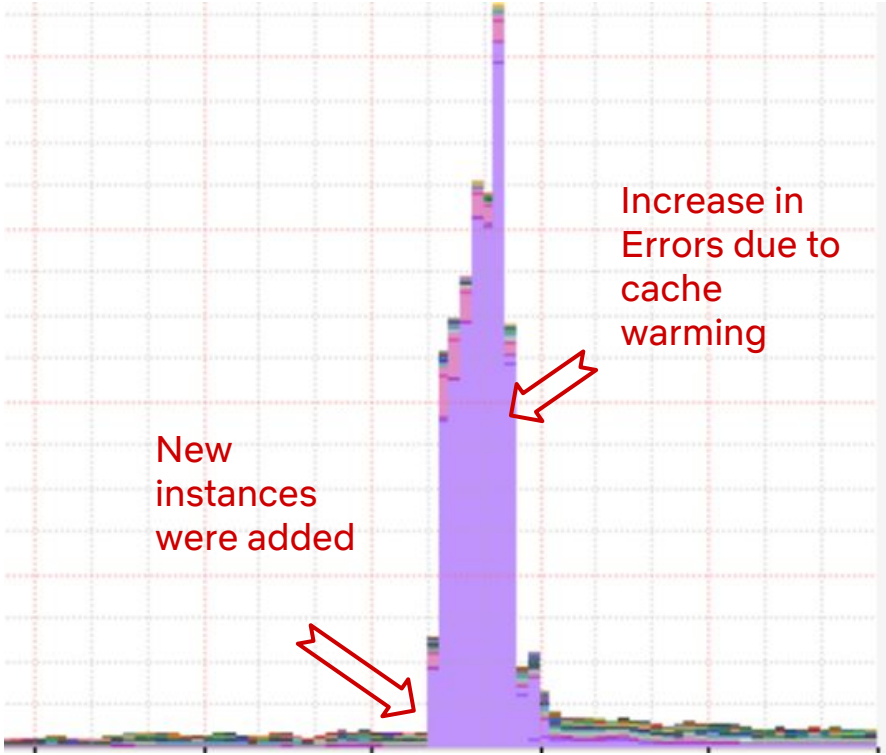

Why Observability over Latency?

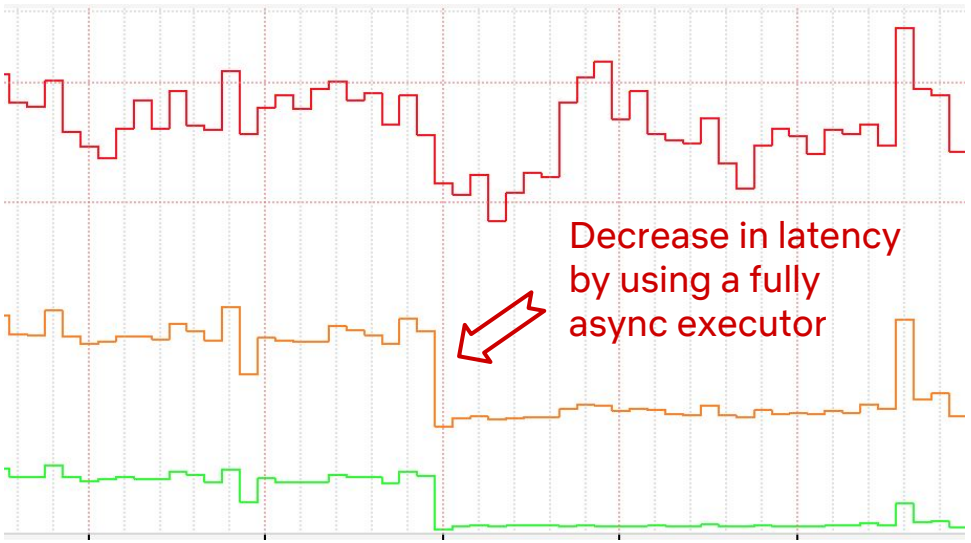

Cost of Async: Loss in Observability

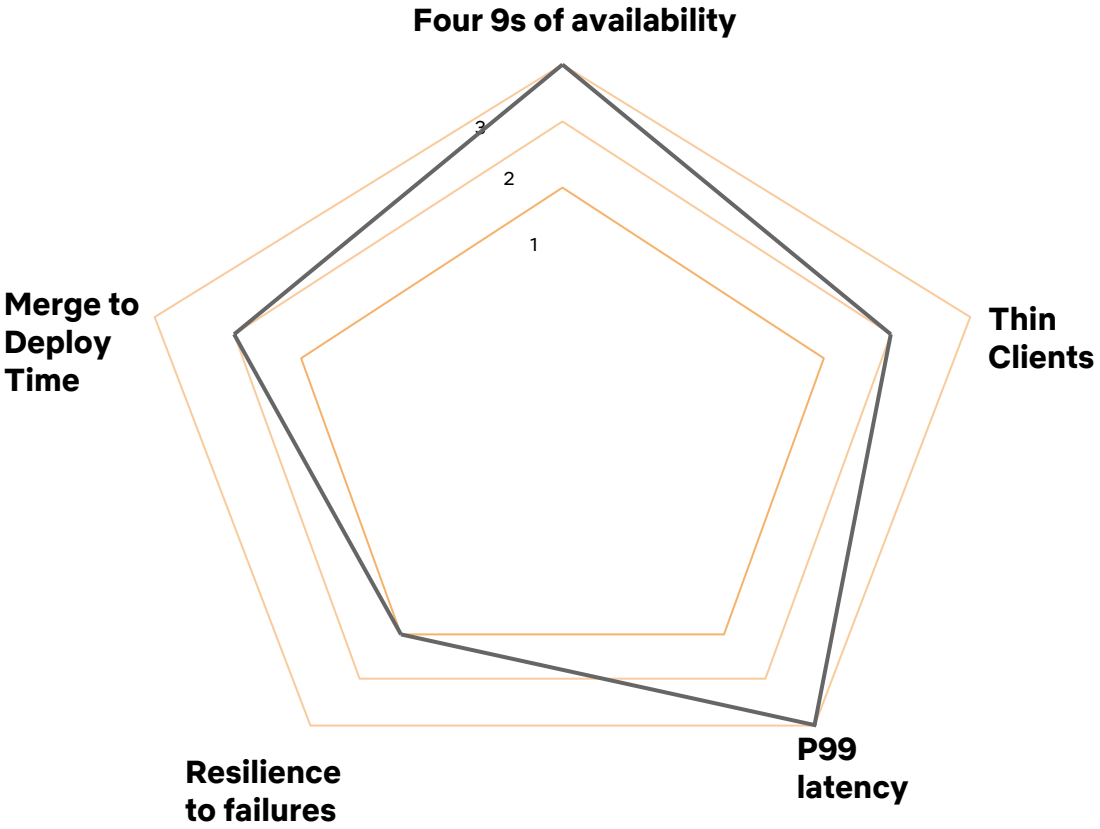
Four 9s of availability
3
2
1
Thin Clients
P99 latency
Resilience to failures
Merge to Deploy Time

Guiding Principle: Define Fitness functions to act as your guide for architectural evolution

Previous Architecture

- Multiple Identities
- Operational Coupling
- Binary Coupling
- Synchronous communication
- Only Java
- Data Monolith

Current Architecture

- Singular Identities
- Operational Isolation
- No Binary Coupling
- Asynchronous communication
- Beyond Java
- Explicit Data Architecture
- Guided Fitness Functions

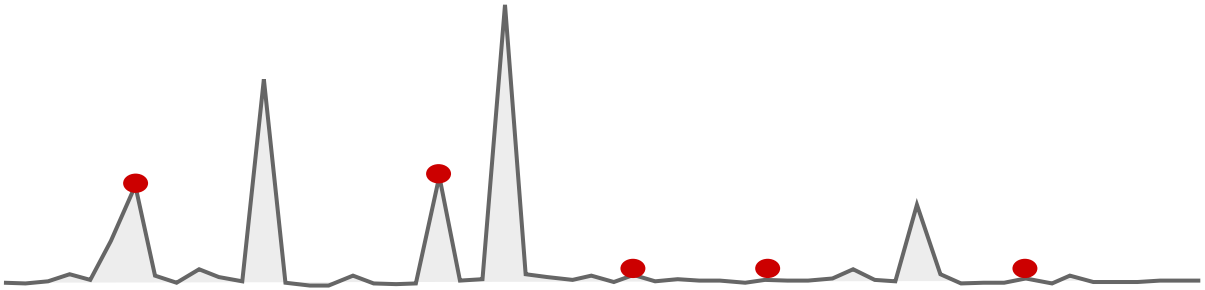

Rollbacks per month Deployments per week

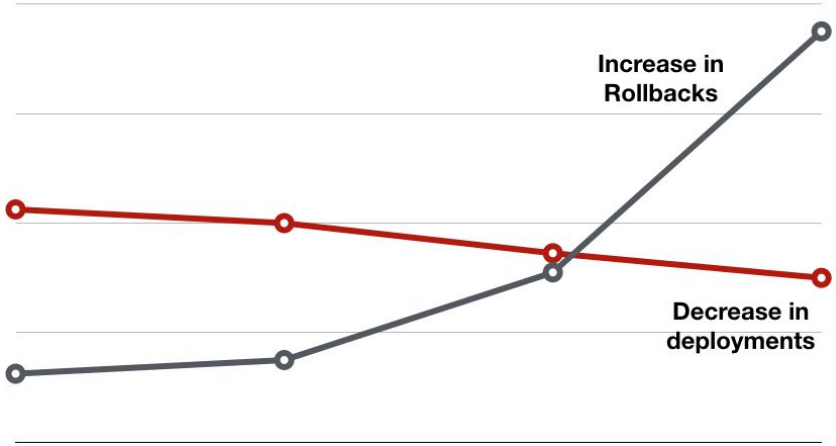

- No incidents in a year
- 4.5 deployments per week
- Just two rollbacks!

Build a Evolutionary Architecture

Identity

Type 1/2 Decisions

Evolvability

NETFLIX